

## Developer Workshop Series Week II



**What we will  
cover:**

- Keyspaces, Tables, Partitions
- The Art of Data Modelling
- Data Types
- What's NEXT?

# Cassandra Cloud-Native Workshop Series

## Developer Workshop Series Week II



**What we will  
cover:**


- Keyspaces, Tables, Partitions
- The Art of Data Modelling
- Data Types
- What's NEXT?

# Cassandra Cloud-Native Workshop Series

# Infrastructure: a Node




A single bare-metal server, a virtual instance or a docker container.



# Infrastructure: a Datacenter (Ring)



A group of nodes located in the same physical location, a cloud datacenter or an availability zone.



# Infrastructure: a Cluster



A group of  
datacenters configured to  
work together.

# Data Structure: a Cell



An intersect of a row and  
a column, stores data.

John
------

# Data Structure: a Row



A single, structured data item in a table.

1	John	Doe	Wizardry
---	------	-----	----------

# Data Structure: a Partition



A group of rows having the same partition token, a base unit of access in Cassandra.

IMPORTANT: stored together, all the rows are guaranteed to be neighbours.

ID	First Name	Last Name	Department
1	John	Doe	Wizardry
399	Marisha	Chapez	Wizardry
415	Maximus	Flavius	Wizardry



# Data Structure: a Table



A group of columns and rows storing partitions.

ID	First Name	Last Name	Department
1	John	Doe	Wizardry
2	Mary	Smith	Dark Magic
3	Patrick	McFadin	DevRel

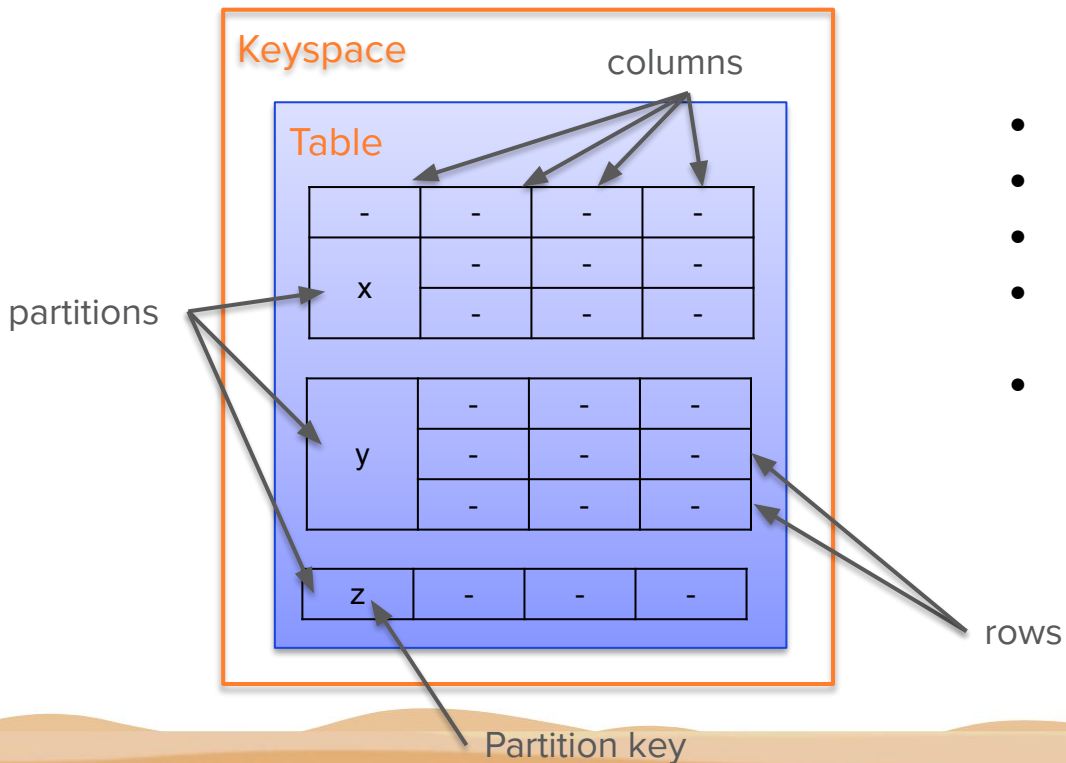
# Data Structure: a Keyspace



A group of tables sharing replication strategy, replication factor and other properties

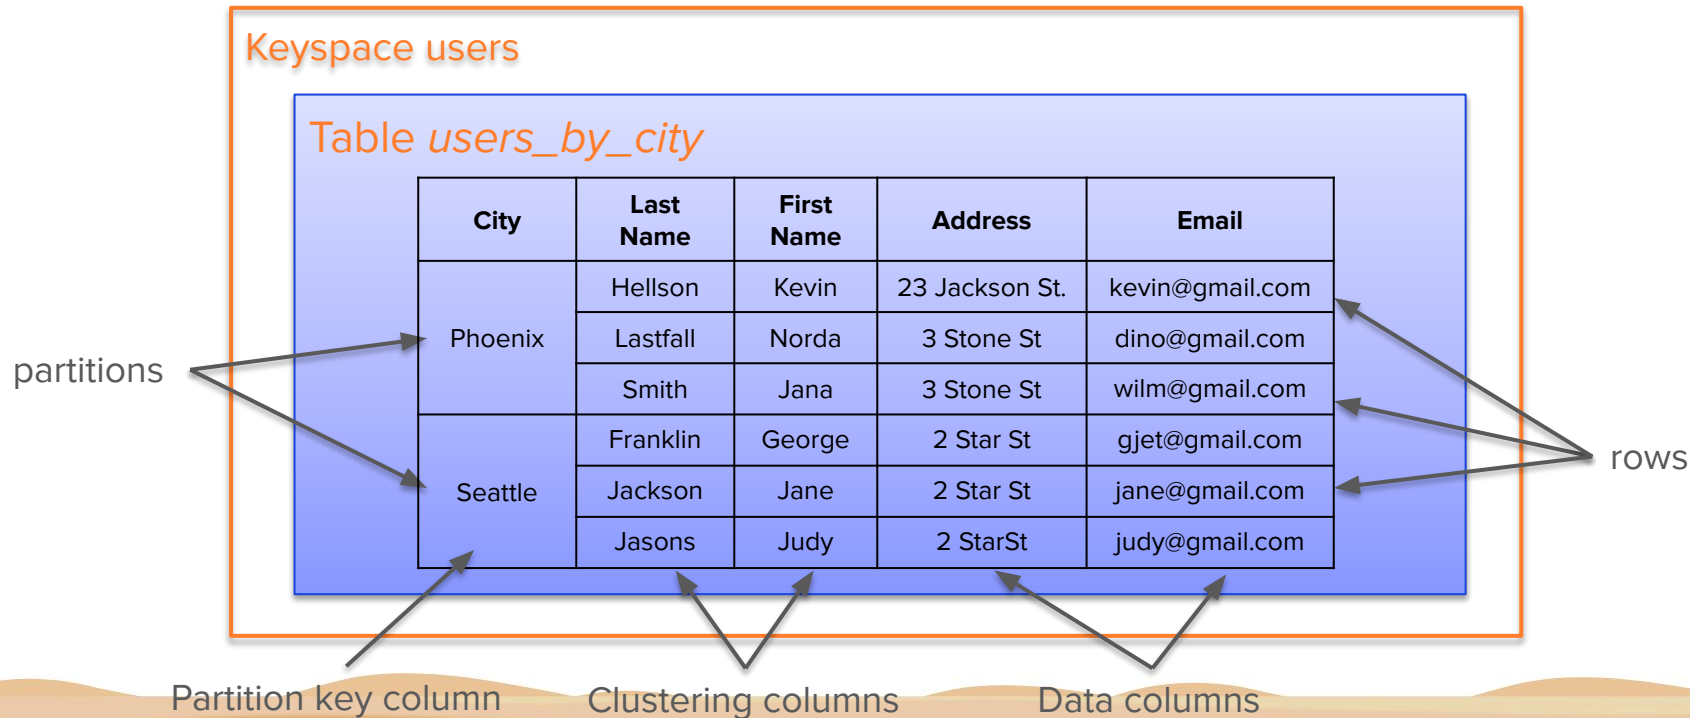
ID	First Name	Last Name	Department
1	John	Doe	Wizardry
2	Mary	Smith	Dark Magic
3	Patrick	McFadin	DevRel

# Data Structure: Overall



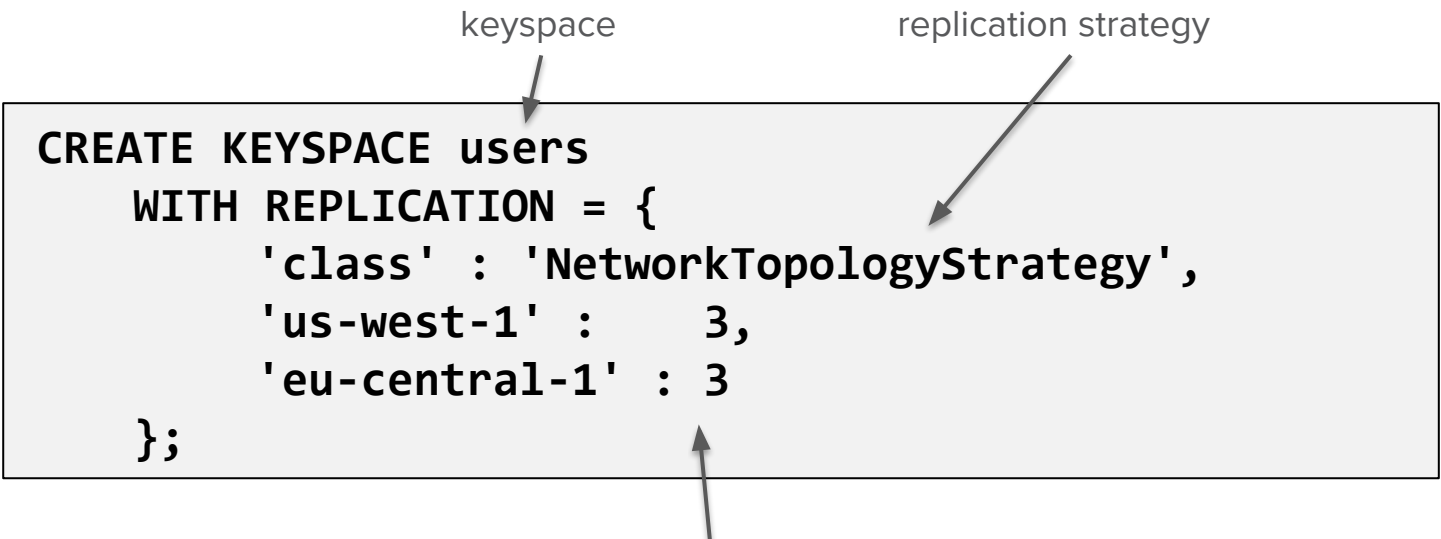
- Tabular data model, with one twist
- *Keyspaces* contain *tables*
- *Tables* are organized in *rows* and *columns*
- Groups of related rows called *partitions* are stored together on the same node (or nodes)
- Each row contains a *partition key*
  - One or more columns that are hashed to determine which node(s) store that data

# Example Data: Users organized by city



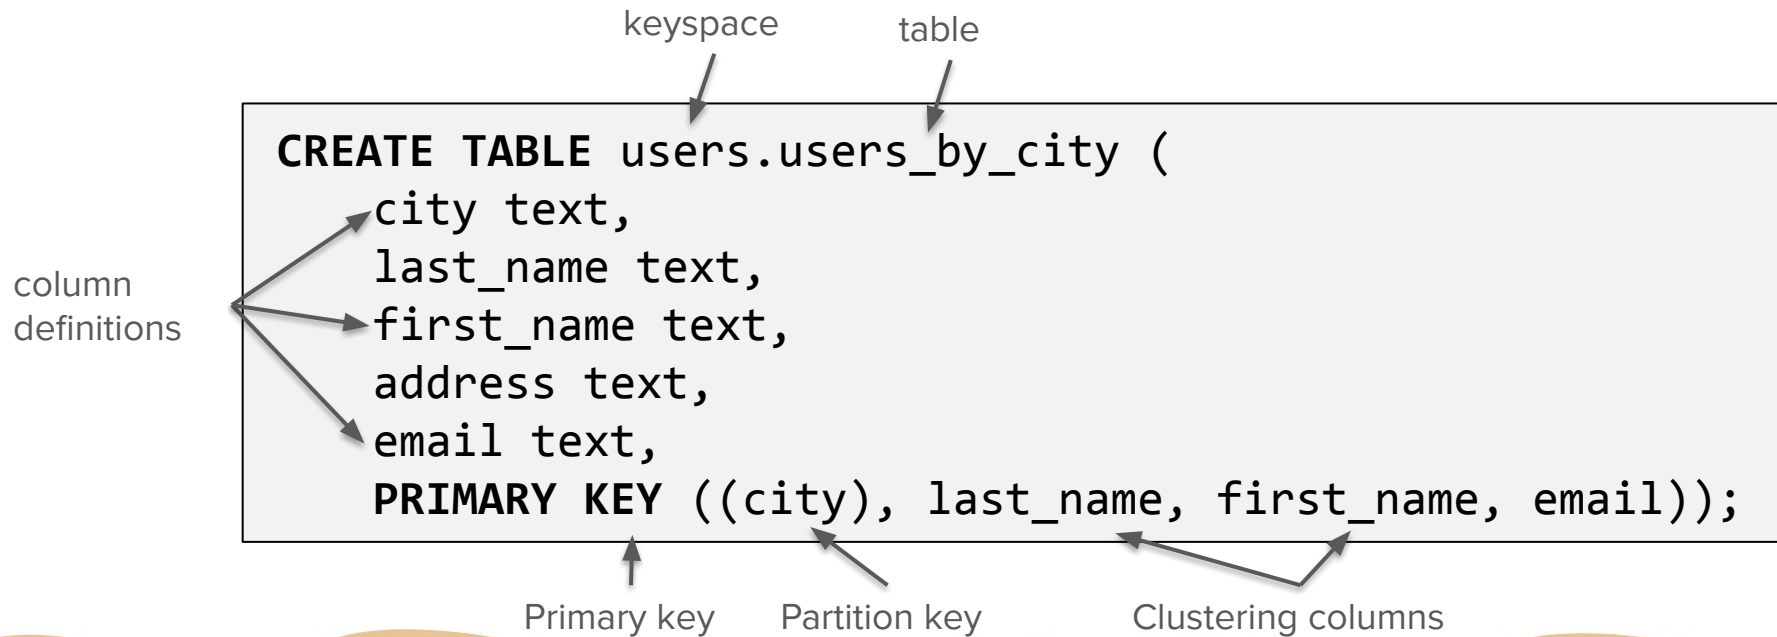
# Creating a Keyspace in CQL

```
CREATE KEYSPACE users
  WITH REPLICATION = {
    'class' : 'NetworkTopologyStrategy',
    'us-west-1' : 3,
    'eu-central-1' : 3
  };
```



Replication factor by data center

# Creating a Table in CQL



# Primary Key

An identifier for a row. Consists of at least one Partition Key and zero or more Clustering Columns.

**MUST ENSURE UNIQUENESS.  
MAY DEFINE SORTING.**

Good Examples:

```
PRIMARY KEY ((city), last_name, first_name, email);
```

```
PRIMARY KEY (user_id);
```

Bad Example:

```
PRIMARY KEY ((city), last_name, first_name);
```

```
CREATE TABLE users.users_by_city (  
  city text,  
  last_name text,  
  first_name text,  
  address text,  
  email text,  
  PRIMARY KEY ((city), last_name, first_name, email));
```

Partition key

Clustering columns

# Partition Key

An identifier for a partition.  
Consists of at least one column,  
may have more if needed

## PARTITIONS ROWS.

Good Examples:

```
PRIMARY KEY (user_id);
```

```
PRIMARY KEY ((video_id), comment_id);
```

Bad Example:

```
PRIMARY KEY ((sensor_id), logged_at);
```

```
CREATE TABLE users.users_by_city (  
  city text,  
  last_name text,  
  first_name text,  
  address text,  
  email text,  
  PRIMARY KEY ((city), last_name, first_name, email));
```

Partition key

Clustering columns



# Clustering Column(s)

Used to ensure uniqueness and sorting order. Optional.

```
CREATE TABLE users.users_by_city (  
  city text,  
  last_name text,  
  first_name text,  
  address text,  
  email text,  
  PRIMARY KEY ((city), last_name, first_name, email));
```

Partition key

Clustering columns

```
PRIMARY KEY ((city), last_name, first_name);
```



Not Unique

```
PRIMARY KEY ((city), last_name, first_name, email);
```



```
PRIMARY KEY ((video_id), comment_id);
```




Not Sorted

```
PRIMARY KEY ((video_id), created_at, comment_id);
```



# Partition: The Beginning

```
CREATE TABLE users.users_by_city (  
  city text,  
  last_name text,  
  first_name text,  
  address text,  
  email text,  
  PRIMARY KEY ((city), last_name, first_name, email));
```



- Every node is responsible for a range of tokens (0-100500, 100501-201000...)
- INSERT a new row, we get the value of its Partition Key (can't be null!)
- We hash this value using MurMur3 hasher <http://murmurhash.shorelabs.com/>  
"Seattle" becomes 2466717130 **Partition Key** = Seattle, **Partition Token** = 2466717130
- This partition belongs to the node[s] responsible for this token
- The INSERT query goes to the nodes storing this partition (Notice Replication Factor)

# Rules of a Good Partition

The Slide of the Year Award!

- Store together what you retrieve together
- Avoid big partitions
- Avoid hot partitions

**PRIMARY KEY** (user\_id);



**PRIMARY KEY** ((video\_id), comment\_id);



**PRIMARY KEY** ((country), user\_id);



# Rules of a Good Partition

The Slide of the Year Award!

- **Store together what you retrieve together**
- Avoid big partitions
- Avoid hot partitions

Example: open a video? Get the comments in a single query!

```
PRIMARY KEY ((video_id), created_at, comment_id);
```



```
PRIMARY KEY ((comment_id), created_at);
```



# Rules of a Good Partition

The Slide of the Year Award!

- Store together what you retrieve together
- **Avoid big partitions**
- Avoid hot partitions

```
PRIMARY KEY ((video_id), created_at, comment_id);
```



```
PRIMARY KEY ((country), user_id);
```



- No technical limitations, but...
- Up to ~100k rows in a partition
- Up to ~100MB in a Partition

# Rules of a Good Partition

The Slide of the Year Award!

- Store together what you retrieve together
- **Avoid big partitions?**
- Avoid hot partitions

Example: a huge IoT infrastructure, hardware all over the world, different sensors reporting their state every 10 seconds. Every sensor reports its UUID, timestamp of the report, sensor's value.

- Sensor ID: UUID
- Timestamp: Timestamp
- Value: float

**PRIMARY KEY** ((sensor\_id), reported\_at);



# Rules of a Good Partition

The Slide of the Year Award!

- Store together what you retrieve together
- **Avoid big and constantly growing partitions!**
- Avoid hot partitions

Example: a huge IoT infrastructure, hardware all over the world, different sensors reporting their state every 10 seconds. Every sensor reports its UUID, timestamp of the report, sensor's value.

- Sensor ID: UUID
- Timestamp: Timestamp
- Value: float

```
PRIMARY KEY ((sensor_id), reported_at);
```



# Rules of a Good Partition

The Slide of the Year Award!

- Store together what you retrieve together
- **Avoid big and constantly growing partitions!**
- Avoid hot partitions

Example: Home IoT infrastructure hardware all over the world, different sensors reporting their status every 10 seconds. Every sensor reports its UUID, time stamp of the report, sensor's value.

Sensor's UUID  
Time stamp:  
Value float

# BUCKETING

```
PRIMARY KEY ((sensor_id), reported_at);
```



```
PRIMARY KEY ((sensor_id, ____), reported_at);
```





# Rules of a Good Partition

The Slide of the Year Award!

- Store together what you retrieve together
- **Avoid big and constantly growing partitions!**
- Avoid hot partitions

Example: a huge IoT infrastructure, hardware all over the world, different sensors reporting their state every 10 seconds. Every sensor reports its UUID, timestamp of the report, sensor's value.

```
PRIMARY KEY ((sensor_id), reported_at);
```



```
PRIMARY KEY ((sensor_id, ____), reported_at);
```



```
PRIMARY KEY ((sensor_id, month_year), reported_at);
```



## BUCKETING

- Sensor ID: UUID
- **MonthYear**: Integer or String
- Timestamp: Timestamp
- Value: float

# Rules of a Good Partition

The Slide of the Year Award!

- Store together what you retrieve together
- Avoid big partitions
- **Avoid hot partitions**

PRIMARY KEY (user\_id);



PRIMARY KEY ((video\_id), created\_at, comment\_id);




PRIMARY KEY ((country), user\_id);




# Always specify the partition key!

If there is no partition key in a query, which node you will ask?



# Always specify the partition key!

If there is no partition key in a query, which node you will ask?



# Always specify the partition key!


If there is no partition key in a query, which node you will ask?

```
CREATE TABLE users.users_by_city (  
  city text,  
  last_name text,  
  first_name text,  
  address text,  
  email text,  
  PRIMARY KEY ((city), last_name, first_name, email));
```

```
SELECT address FROM users_by_city WHERE first_name = "Anna";
```



```
SELECT address FROM users_by_city WHERE city = "Otterberg" AND last_name = "Koshkina";
```



## Developer Workshop Series Week II





**What we will  
cover:**

- Keyspaces, Tables, Partitions
- The Art of Data Modelling
- Data Types
- What's NEXT?

# Cassandra Cloud-Native Workshop Series

# Relational Data Modelling

1. Analyze raw data
2. Identify entities, their properties and relations
3. Design tables, using normalization and foreign keys.
4. Use JOIN when doing queries to join denormalized data from multiple tables



Employees

userId	firstName	lastName
1	Edgar	Codd
2	Raymond	Boyce

Department

departmentId	department
1	Engineering
2	Math

Arrows indicate a relationship between the 'Employees' table and the 'Department' table, specifically linking the 'userId' column to the 'departmentId' column.

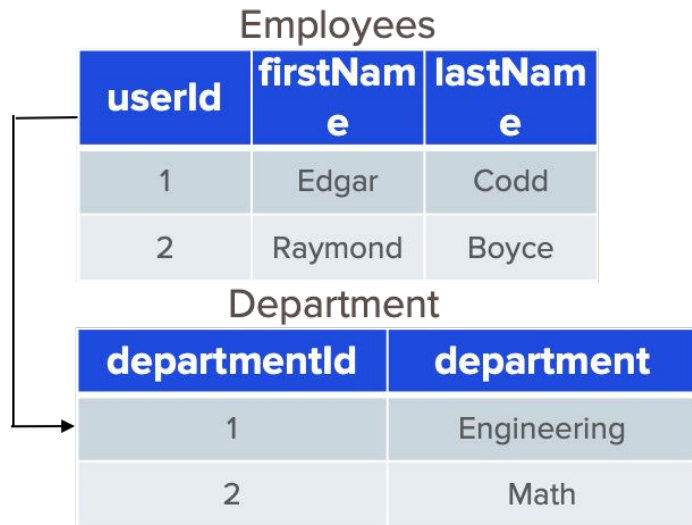


# Normalization

“Database normalization is the process of structuring a relational database in accordance with a series of so-called normal forms in order to reduce data redundancy and improve data integrity. It was first proposed by Edgar F. Codd as part of his relational model.”

**PROS:** Simple write, Data Integrity

**CONS:** Slow read, Complex Queries





# Denormalization

“Denormalization is a strategy used on a database to increase performance. In computing, denormalization is the process of trying to improve the read performance of a database, at the expense of losing some write performance, by adding redundant copies of data”

**PROS:** Quick Read, Simple Queries

**CONS:** Multiple Writes, Manual Integrity

Employees


userId	firstName	lastName	department
1	Edgar	Codd	Engineering
2	Raymond	Boyce	Math

Department


departmentId	department
1	Engineering
2	Math

# NoSQL Data Modelling


1. Analyze user behaviour (customer first!)
2. Identify workflows, their dependencies and needs
3. Define Queries to fulfill these workflows
4. Knowing the queries, design tables, using denormalization.
5. Use BATCH when inserting or updating denormalized data of multiple tables



id	firstName	lastName	department
1	Edgar	Codd	Engineering
2	Raymond	Boyce	Math



# Let's go practical!

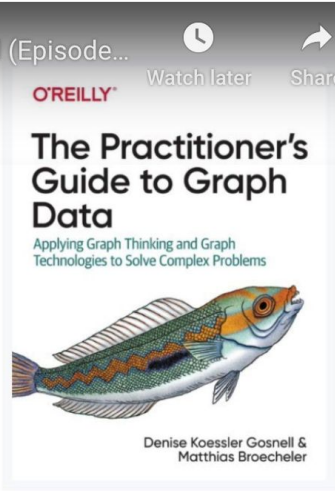


Graph Data and Code with Denise & David (Episode...)

Watch later Share

## Graph Data & Code with Denise and David

Tuesdays at 16:00 GMT/12:00 EDT/9:00 PDT



### DATASTAX

0.0 overall rating


by: [Adaline Walker](#)

Added on 7/8/2020

Join Dr. Denise Gosnell as she continues to school David on their Graph journey from The Practitioner's Guide to Graph Data. This week we will:


Show More

#### LATEST COMMENTS

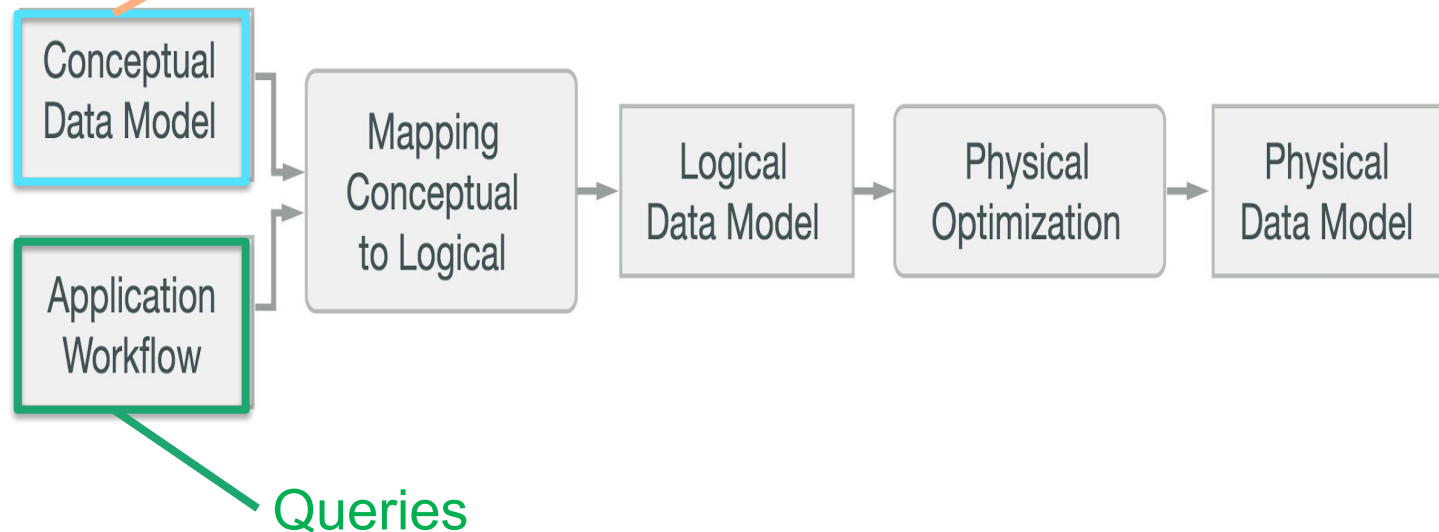
 [Aleksandr Volochnev](#) a few seconds ago  
Great video! Thank you!

Leave a comment

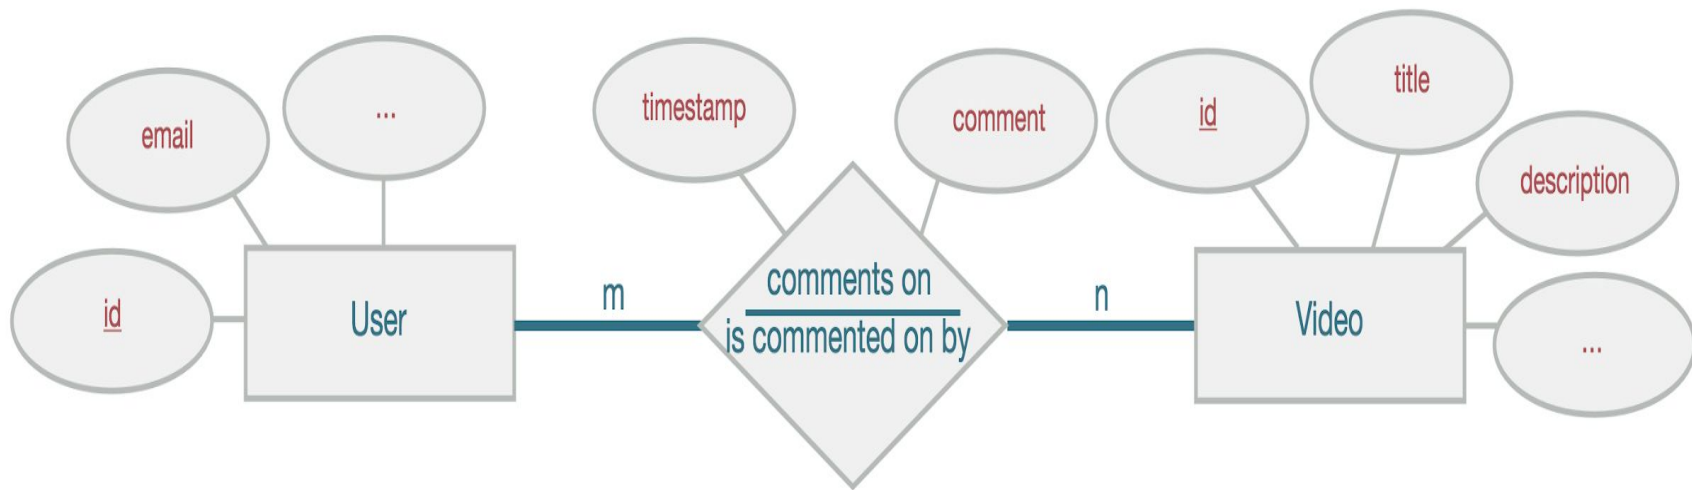

# Designing Process: Step by Step




Entities & Relationships



# Designing Process: Conceptual Data Model



# Designing Process: Application Workflow



## Use-Case I:

- A User opens a Video Page


**WF1:** Find **comments** related to target **video** using its identifier, most recent first

## Use-Case II + III:

- A User opens a Profile
- A Moderator verifies a User if spammer or not

**WF2:** Find **comments** related to target **user** using its identifier, get most recent first

# Designing Process: Mapping



**Query I:** Find comments posted for a user with a known id (show most recent first)




comments\_by\_user

**Query II:** Find comments for a video with a known id (show most recent first)




comments\_by\_video


# Designing Process: Mapping



```
SELECT * FROM comments_by_user  
WHERE userid = <some UUID>
```




```
SELECT * FROM comments_by_video  
WHERE videoid = <some UUID>
```





# Designing Process: Logical Data Model




## comments\_by\_user

userid	K
creationdate	C ↓
commentid	C ↑
videoid	
comment	

## comments\_by\_video

videoid	K
creationdate	C ↓
commentid	C ↑
userid	
comment	


# Designing Process: Physical Data Model



comments_by_user			
userid	UUID	K	
commentid	TIMEUUID	C	↓
videoid	UUID		
comment	TEXT		

comments_by_video			
videoid	UUID	K	
commentid	TIMEUUID	C	↓
userid	UUID		
comment	TEXT		

# Designing Process: Schema DDL



```
CREATE TABLE IF NOT EXISTS comments_by_user (  
    userid uuid,  
    commentid timeuuid,  
    videoid uuid,  
    comment text,  
    PRIMARY KEY ((userid), commentid)  
) WITH CLUSTERING ORDER BY (commentid DESC);
```

```
CREATE TABLE IF NOT EXISTS comments_by_video (  
    videoid uuid,  
    commentid timeuuid,  
    userid uuid,  
    comment text,  
    PRIMARY KEY ((videoid), commentid)  
) WITH CLUSTERING ORDER BY (commentid DESC);
```

# menti.com

89 49 47



## Developer Workshop Series Week II



**What we will  
cover:**

- Keyspaces, Tables, Partitions
- The Art of Data Modelling
- Data Types
- What's NEXT?


# Cassandra Cloud-Native Workshop Series

# Basic Data Types

type	constants supported	description
<code>ascii</code>	<code>string</code>	ASCII character string
<code>bigint</code>	<code>integer</code>	64-bit signed long
<code>blob</code>	<code>blob</code>	Arbitrary bytes (no validation)
<code>boolean</code>	<code>boolean</code>	Either <code>true</code> or <code>false</code>
<code>counter</code>	<code>integer</code>	Counter column (64-bit signed value). See <a href="#">Counters</a> for details
<code>date</code>	<code>integer</code> , <code>string</code>	A date (with no corresponding time value). See <a href="#">Working with dates</a> below for details
<code>decimal</code>	<code>integer</code> , <code>float</code>	Variable-precision decimal
<code>double</code>	<code>integer</code> <code>float</code>	64-bit IEEE-754 floating point
<code>duration</code>	<code>duration</code> ,	A duration with nanosecond precision. See <a href="#">Working with durations</a> below for details
<code>float</code>	<code>integer</code> , <code>float</code>	32-bit IEEE-754 floating point
<code>inet</code>	<code>string</code>	An IP address, either IPv4 (4 bytes long) or IPv6 (16 bytes long). Note that there is no <code>inet</code> constant, IP address should be input as strings
<code>int</code>	<code>integer</code>	32-bit signed int
<code>smallint</code>	<code>integer</code>	16-bit signed int
<code>text</code>	<code>string</code>	UTF8 encoded string
<code>time</code>	<code>integer</code> , <code>string</code>	A time (with no corresponding date value) with nanosecond precision. See <a href="#">Working with times</a> below for details
<code>timestamp</code>	<code>integer</code> , <code>string</code>	A timestamp (date and time) with millisecond precision. See <a href="#">Working with timestamps</a> below for details
<code>timeuuid</code>	<code>uuid</code>	Version 1 <code>UUID</code> , generally used as a “conflict-free” timestamp. Also see <a href="#">Timeuuid functions</a>
<code>tinyint</code>	<code>integer</code>	8-bit signed int
<code>uuid</code>	<code>uuid</code>	A <code>UUID</code> (of any version)
<code>varchar</code>	<code>string</code>	UTF8 encoded string
<code>varint</code>	<code>integer</code>	Arbitrary-precision integer

# Collections

I'm a SET with a bunch  
of unordered things



I'm an ordered LIST

0	1	2	3	4	6
---	---	---	---	---	---

I'm a MAP of  
key/value pairs

Key	Value
K1	V1
K2	V2
K3	V3
K4	V4
K5	V5

# Collection: Set

```
CREATE killrvideo.videos (  
  videoid          uuid,  
  userid          uuid,  
  name            text,  
  description      text,  
  location        text,  
  location_type    int,  
  preview_image_location text,  
  tags            set<text>,  
  added_date      timestamp,  
  PRIMARY KEY (videoid)  
);
```

{‘Family’, ‘Disney’, ‘Princess’}

{‘Thriller’, ‘Short’}

{‘Tragicomedy’, ‘Western’}



## Collection: Set

```
INSERT INTO killrvideo.videos (videoid, tags)
VALUES(12345678-1234-1234-1234-123456789012,
{'Side-splitter', 'Short'});
```

Insert

```
UPDATE killrvideo.videos
SET tags = {'Dark', 'Sad'}
WHERE videoid = 12345678-1234-1234-1234-123456789012;
```

Replace entire set

```
UPDATE killrvideo.videos
SET tags = tags + {'Enthralling'}
WHERE videoid = 12345678-1234-1234-1234-123456789012;
```

Add to set

## Collection: List

```
CREATE killrvideo.actors_by_video (  
    videoid uuid,  
    actors list<text>, // alphabetical list of actors  
    PRIMARY KEY (videoid)  
);
```



## Collection: List

```
INSERT INTO killrvideo.actors_by_video (videoid, actors)  
VALUES(12345678-1234-1234-1234-123456789012,  
['Adams', 'Baker', 'Cox']);
```

Insert

```
UPDATE killrvideo.actors_by_video  
SET actors = ['Arthur', 'Beverly']  
WHERE videoid = 12345678-1234-1234-1234-123456789012;
```

Replace entire list

```
UPDATE killrvideo.actors_by_video  
SET actors=actors + ['Crawford']  
WHERE videoid = 12345678-1234-1234-1234-123456789012;
```

Append

## Collection: List


```
UPDATE killrvideo.actors_by_video  
  SET actors[1] = 'Brown'  
WHERE videoid = 12345678-1234-1234-1234-123456789012;
```

Replace an element

Note: replacing an element requires a read-before-write, which implies performance penalty.

## Collection: Map

```
CREATE TABLE killrvideo.users(  
  userid      uuid,  
  phone_nos   map<text, text>,  
  PRIMARY KEY (userid)  
);
```





## Collection: Map

```
INSERT INTO killrvideo.users (userid, phone_nos)
VALUES(12345678-1234-1234-1234-123456789012,
{'cell':'867-5309', 'home':'555-1212',
'busi':'800-555-1212'});
```

Insert

```
UPDATE killrvideo.users
SET phone_nos = {'cell':'867-5310', 'office':'555-1212'}
WHERE userid = 12345678-1234-1234-1234-123456789012;
```

Replace entire map


```
UPDATE killrvideo.users
SET phone_nos = phone_nos + {'desk': '270-555-1213'}
WHERE userid = 12345678-1234-1234-1234-123456789012;
```

Add to map

# User Defined Types

```
CREATE TYPE killrvideo.address(  
  street text,  
  city   text,  
  state  text,  
);
```

```
CREATE TABLE killrvideo.users(  
  userid      uuid,  
  location    address,  
  PRIMARY KEY (userid)  
);
```



# User Defined Types

```
INSERT INTO killrvideo.users (userid, location)  
VALUES(12345678-1234-1234-1234-123456789012,  
{street:'123 Main', city:'Metropolis', state:'CA'});
```

Insert

```
UPDATE killrvideo.users  
SET location = {street:'234 Elm', city:'NYC', state:'NY'}  
WHERE userid = 12345678-1234-1234-1234-123456789012;
```

Replace entire UDT

```
UPDATE killrvideo.users  
SET location.city = 'Albany'  
WHERE userid = 12345678-1234-1234-1234-123456789012;
```

Replace one UDT field



# User Defined Types

```
SELECT location.city FROM killrvideo.users  
WHERE userid = 12345678-1234-1234-1234-123456789012;
```

Select field

# Counters

- 64-bit signed integer
- Use-case:
  - Imprecise values such as likes, views, etc.
- Two operations:
  - Increment
  - Decrement
  - First op assumes the value is zero

# Counters

- Cannot be part of primary key
- Counters not mixed with other types in table
- Value cannot be set
- Rows with counters cannot be inserted
- Updates are not idempotent
  - Counters should *not* be used for precise values

# Counters

```
CREATE TABLE killrvideo.video_playback_stats (  
    videoid uuid,  
    views counter,  
    PRIMARY KEY (videoid)  
);
```

# Counters

Incrementing a counter:

```
UPDATE killrvideo.videos SET views = views + 1  
WHERE videoid = 12345678-1234-1234-1234-123456789012;
```

This format must be observed

This can be an integer value

Decrementing a counter:

```
UPDATE killrvideo.videos SET views = views - 1  
WHERE videoid = 12345678-1234-1234-1234-123456789012;
```

Just change the sign