

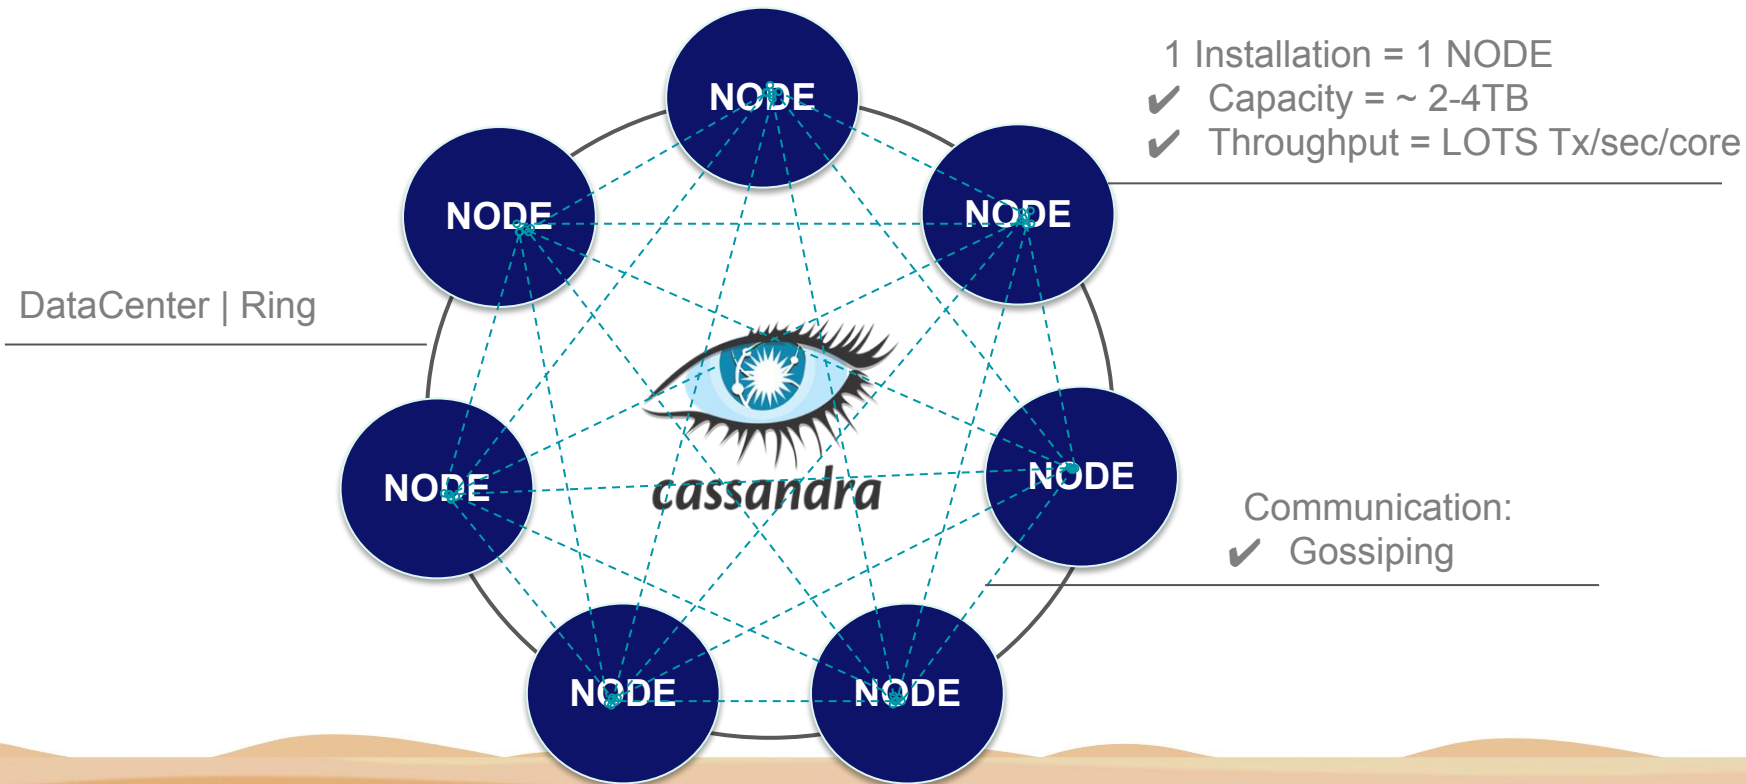
Developer Workshop Series Week 1



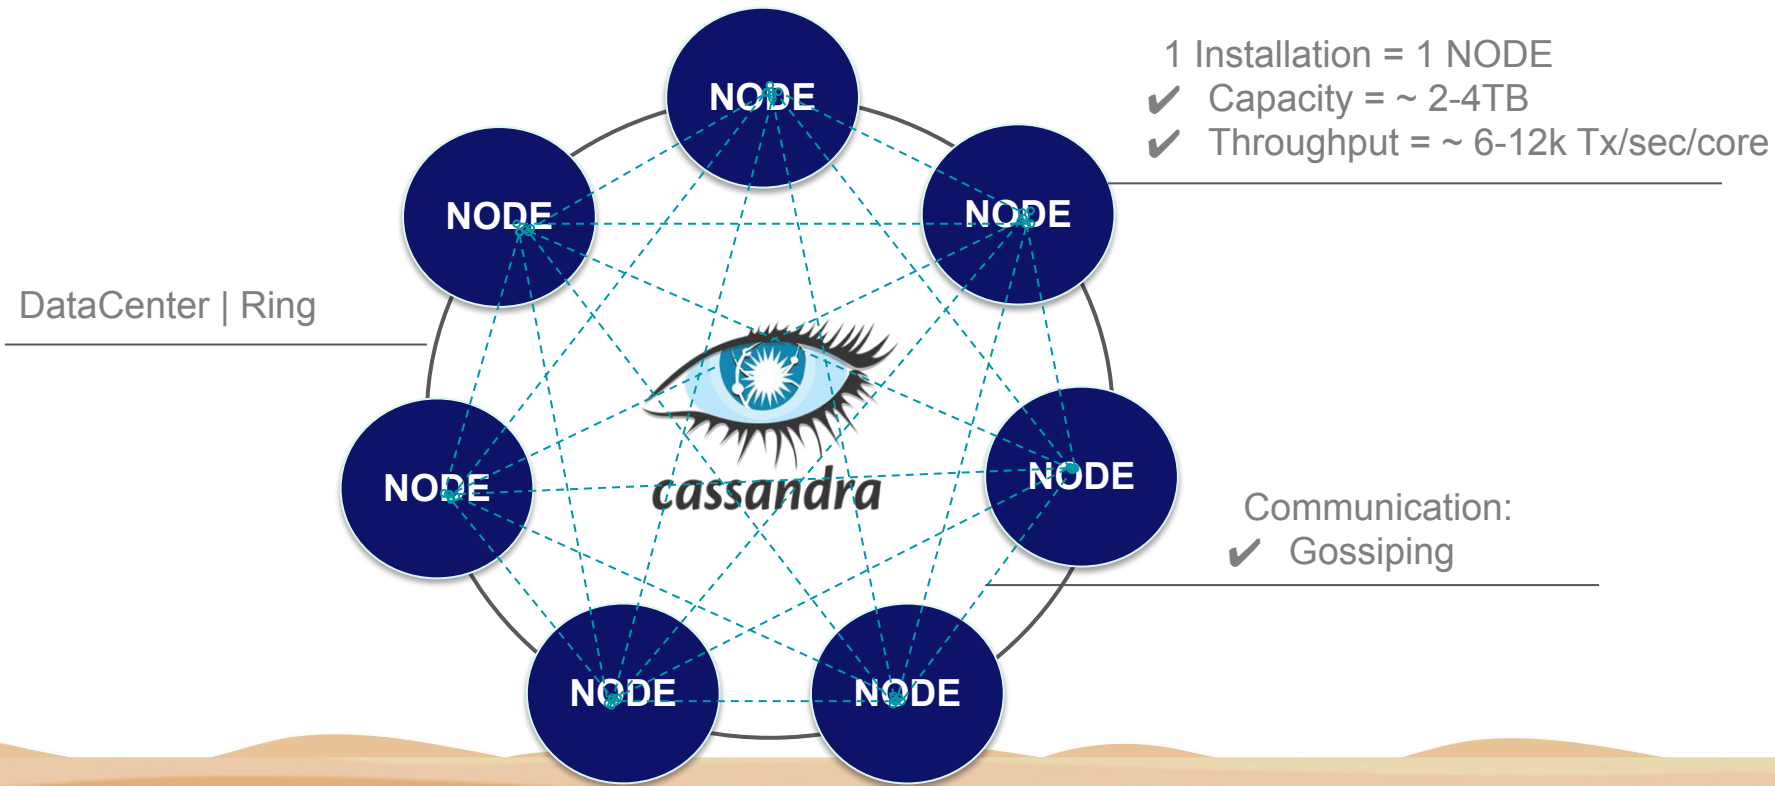
**What we will
cover:**

- Bootstrapping
- Apache Cassandra™ Why, What & When
- Read and Write path
- Uber High Level Data Modeling
- What's NEXT?

Apache Cassandra™ = NoSQL Distributed Database




Apache Cassandra™ = NoSQL Distributed Database



Scales Linearly


- Need more capacity?
- Need more throughput?
- Add nodes!

Balanced Read/Write Mix




Horizontal vs. Vertical Scaling


- Vertical scaling requires one large expensive machine
- Horizontal scaling requires multiple less-expensive commodity hardware



100,000 transactions/second



200,000 transactions/second



400,000 transactions/second

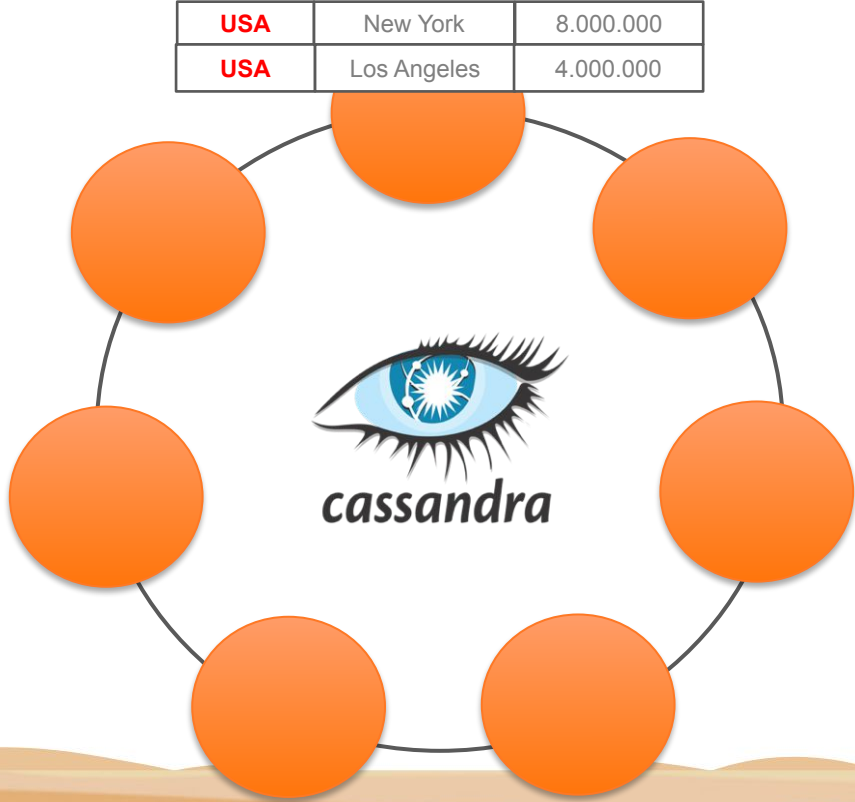
Data is Distributed



Country	City	Population
USA	New York	8.000.000
USA	Los Angeles	4.000.000
FR	Paris	2.230.000
DE	Berlin	3.350.000
UK	London	9.200.000
AU	Sydney	4.900.000
DE	Nuremberg	500.000
CA	Toronto	6.200.000
CA	Montreal	4.200.000
FR	Toulouse	1.100.000
JP	Tokyo	37.430.000
IN	Mumbai	20.200.000

Partition Key

Data is Distributed

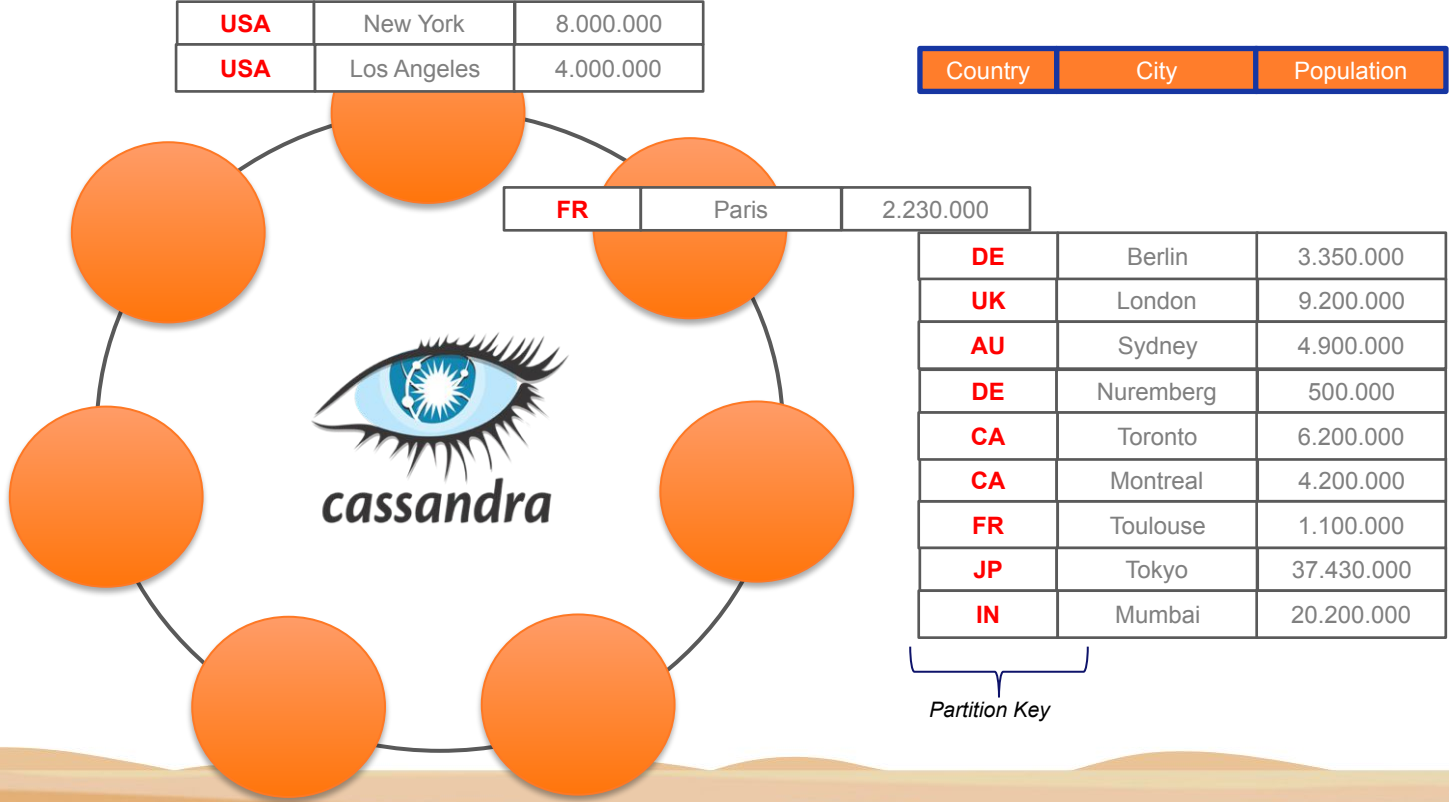


Country	City	Population
---------	------	------------

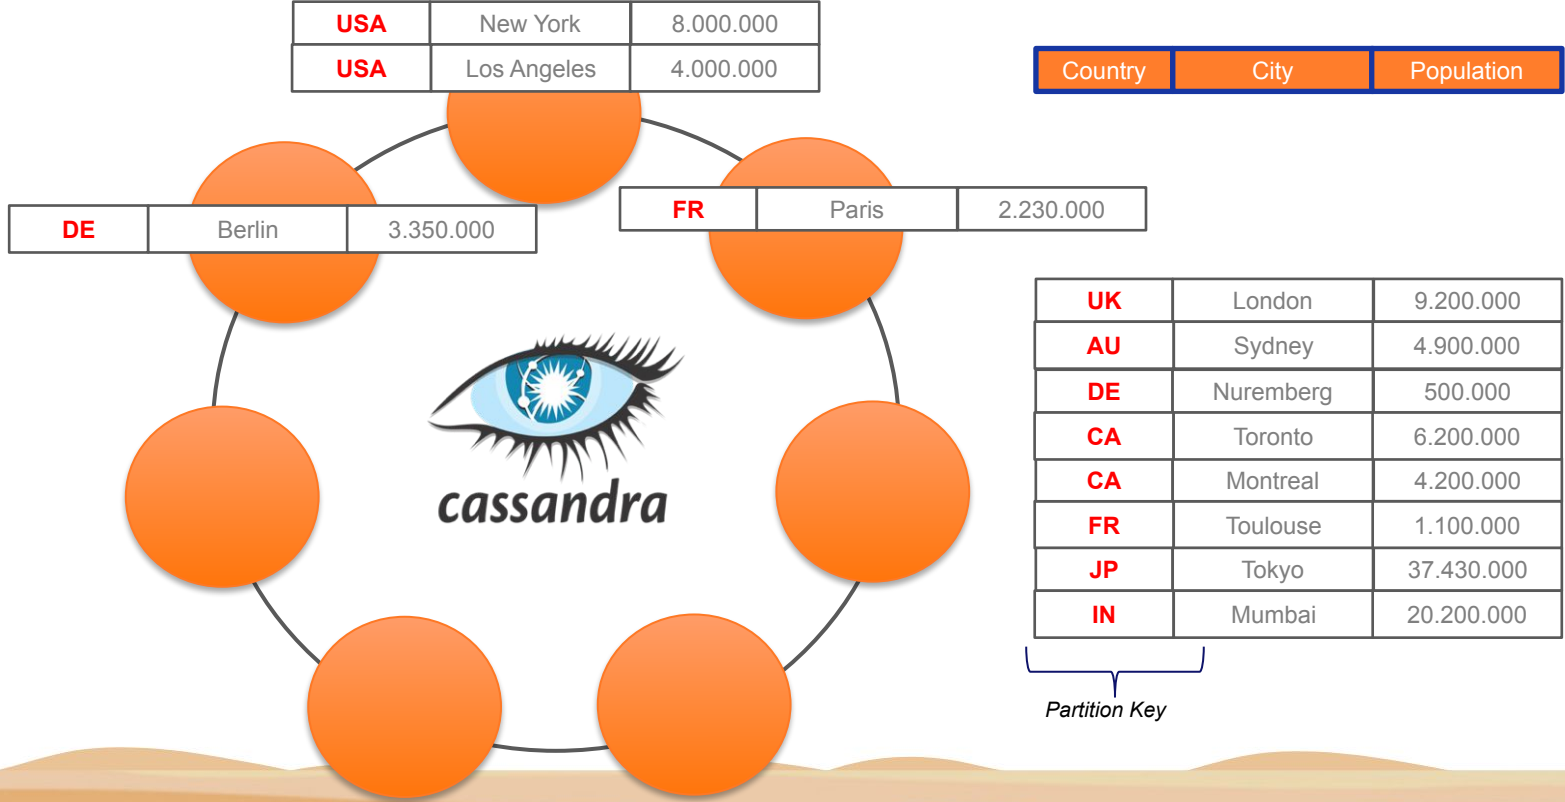
FR	Paris	2.230.000
DE	Berlin	3.350.000
UK	London	9.200.000
AU	Sydney	4.900.000
DE	Nuremberg	500.000
CA	Toronto	6.200.000
CA	Montreal	4.200.000
FR	Toulouse	1.100.000
JP	Tokyo	37.430.000
IN	Mumbai	20.200.000

Partition Key

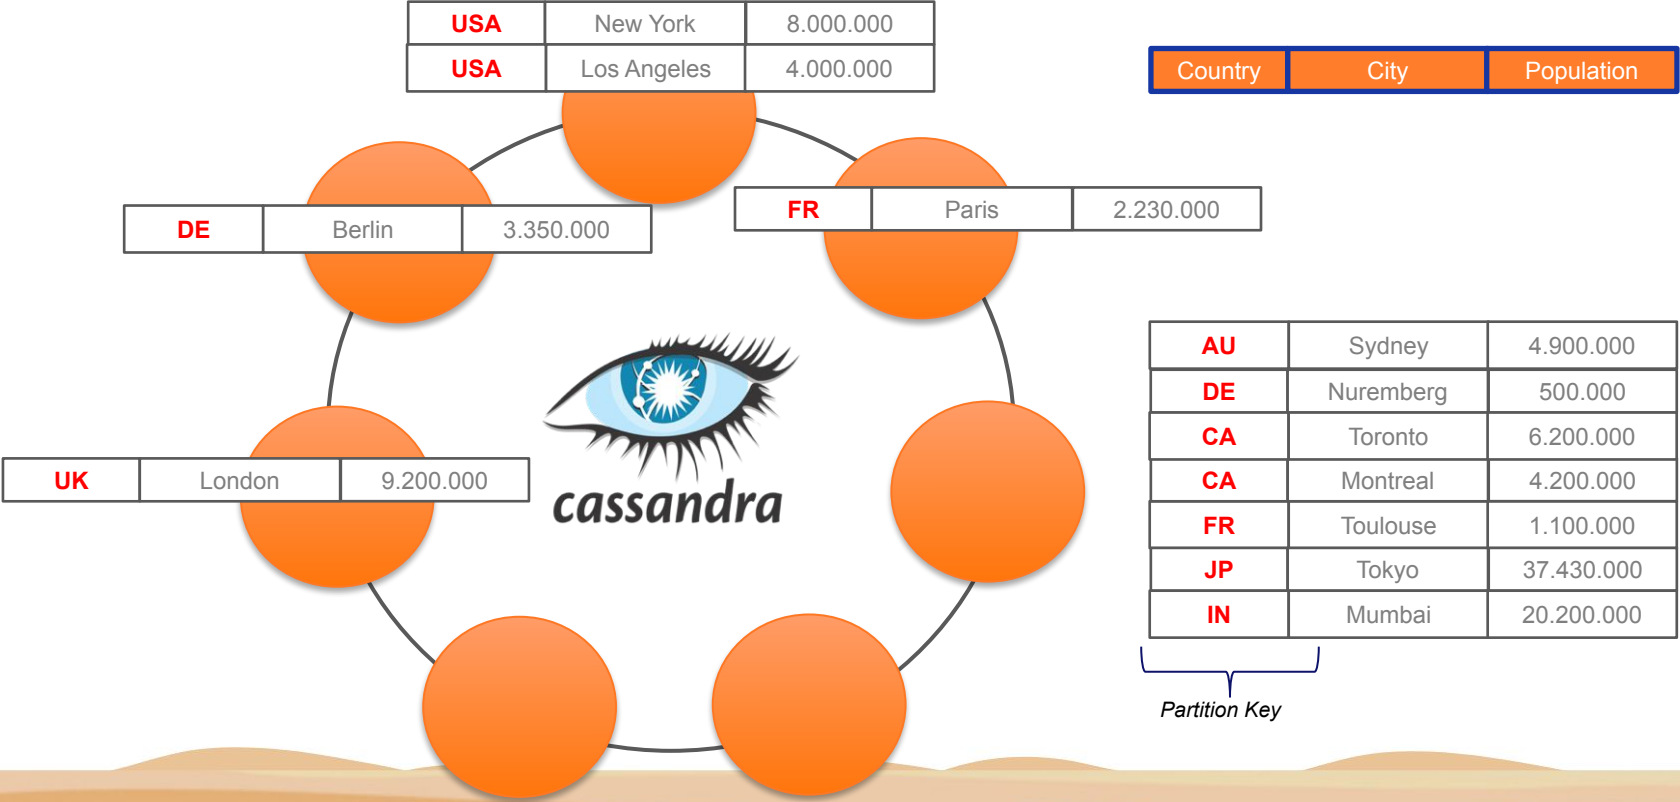
Data is Distributed



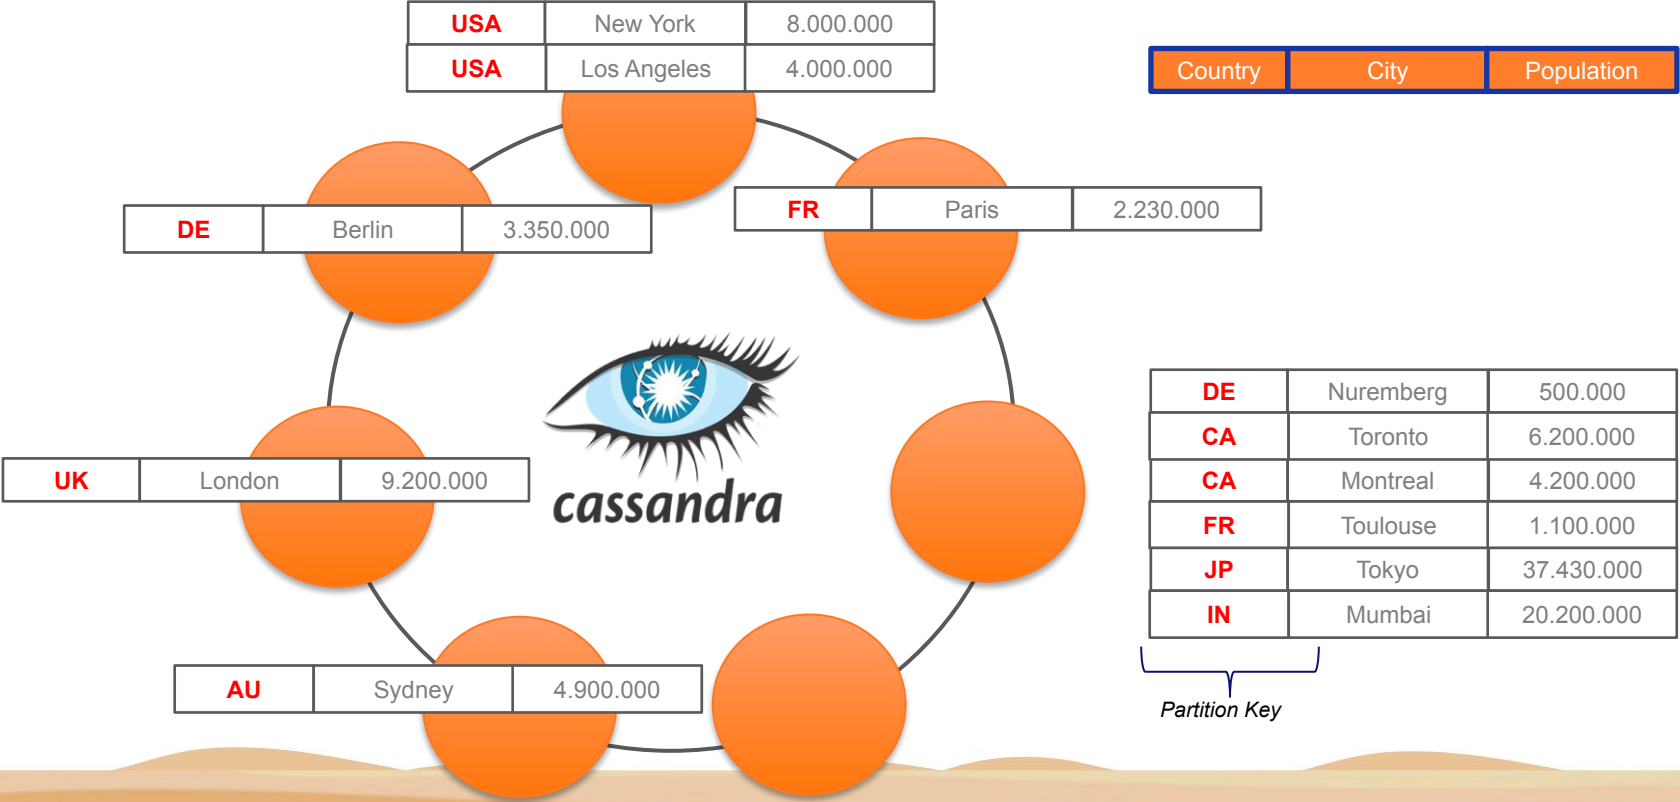
Data is Distributed



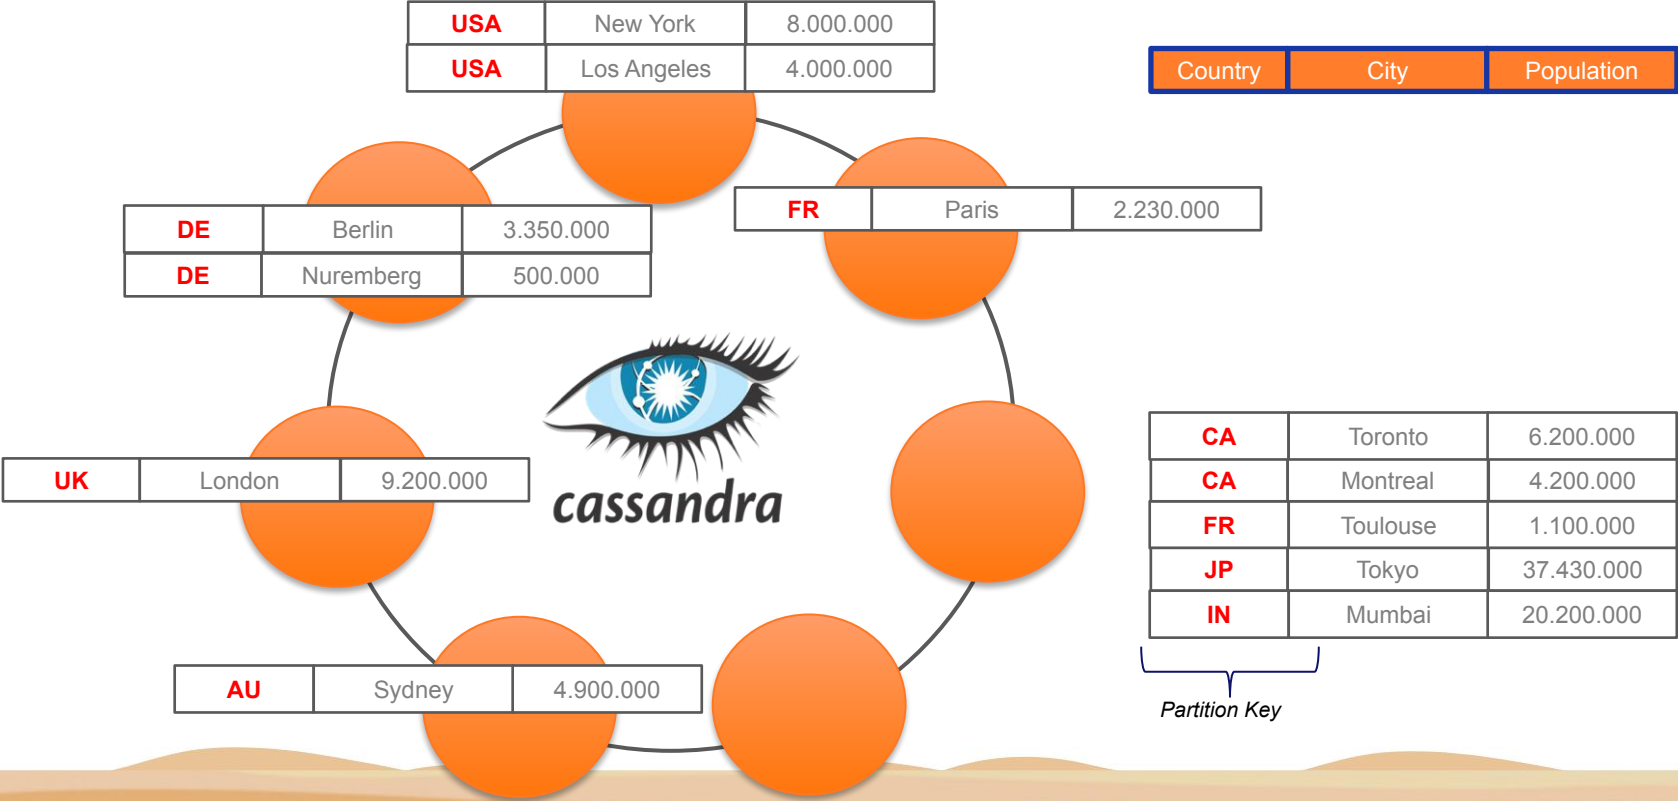
Data is Distributed



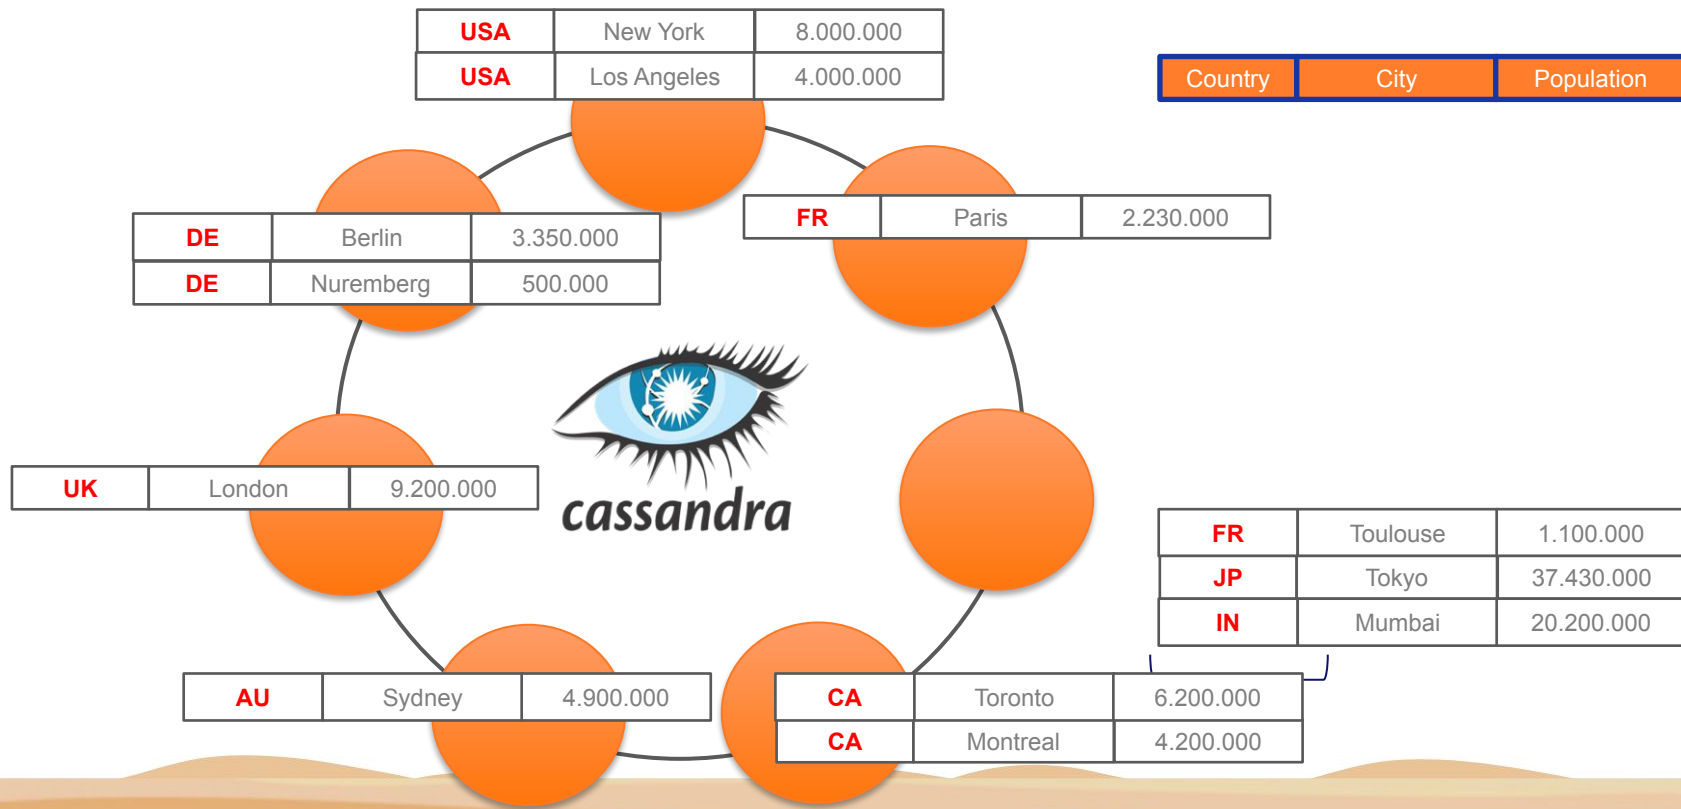
Data is Distributed



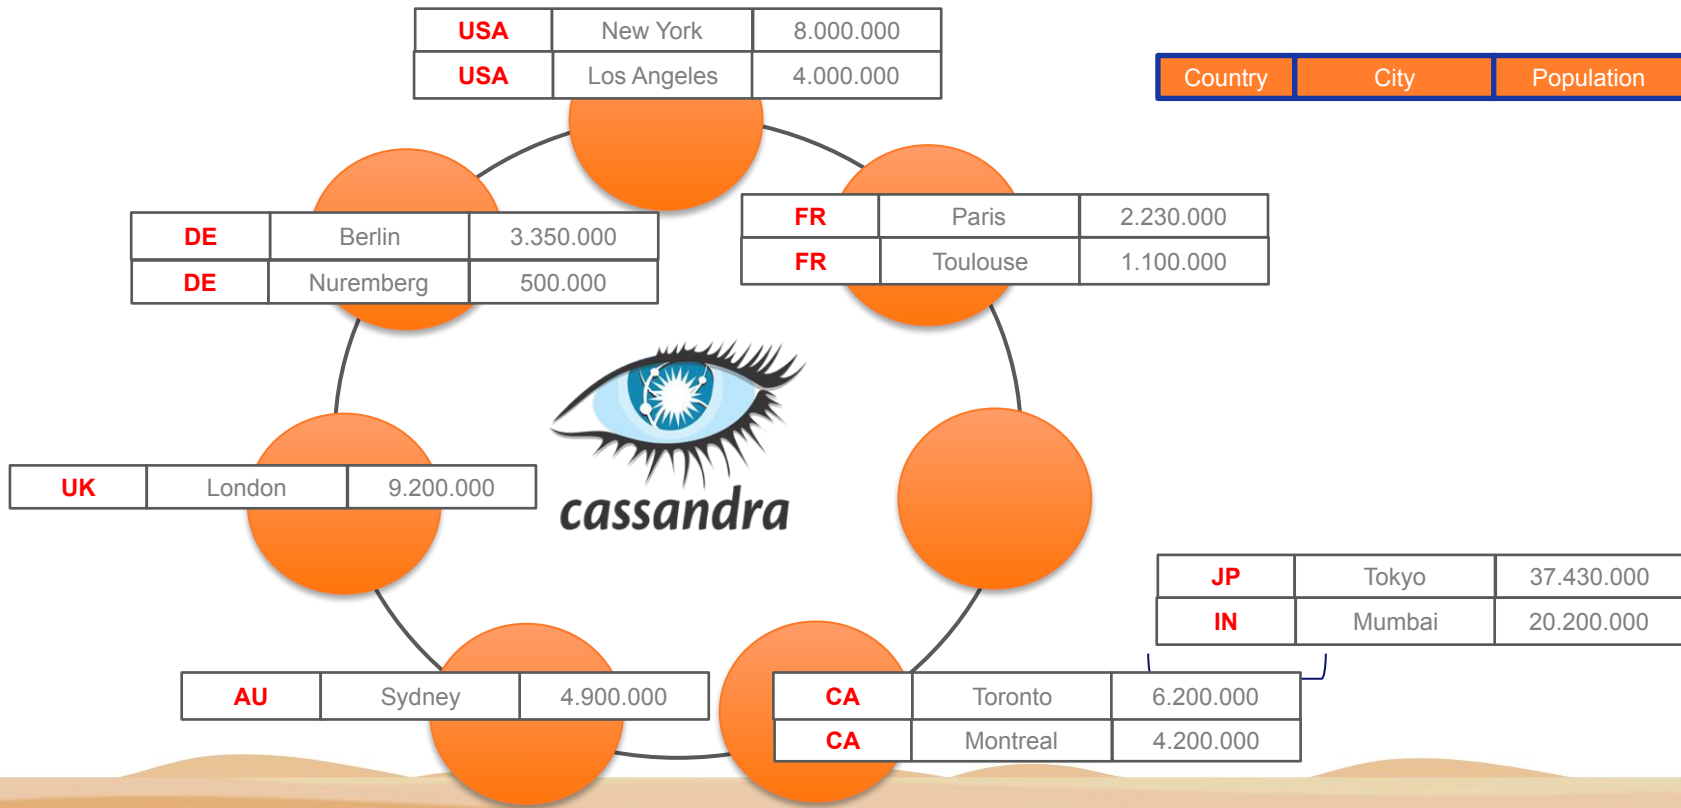
Data is Distributed



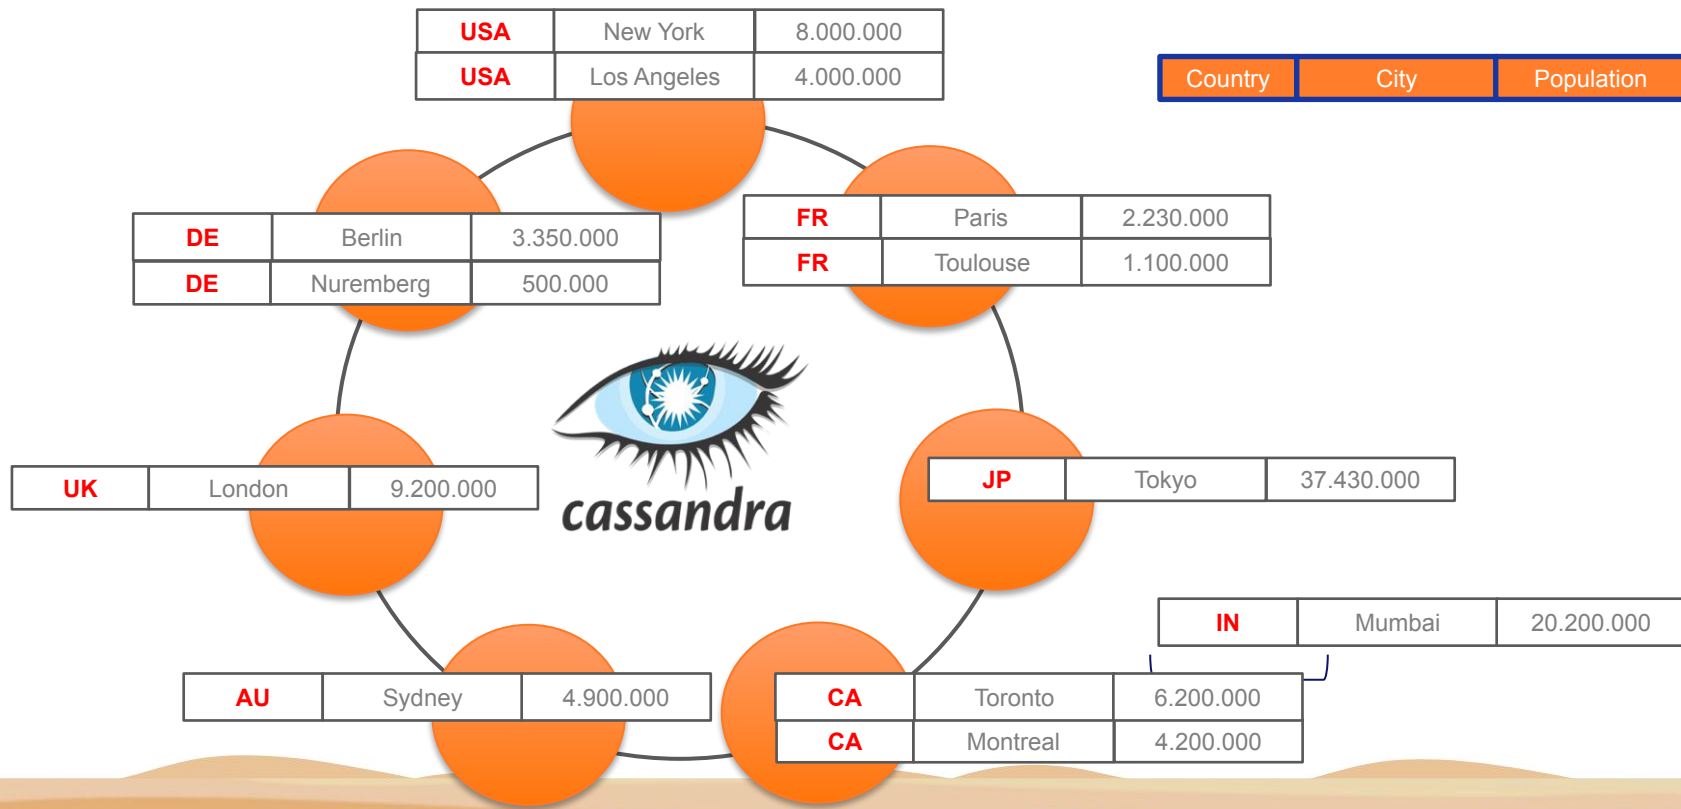
Data is Distributed



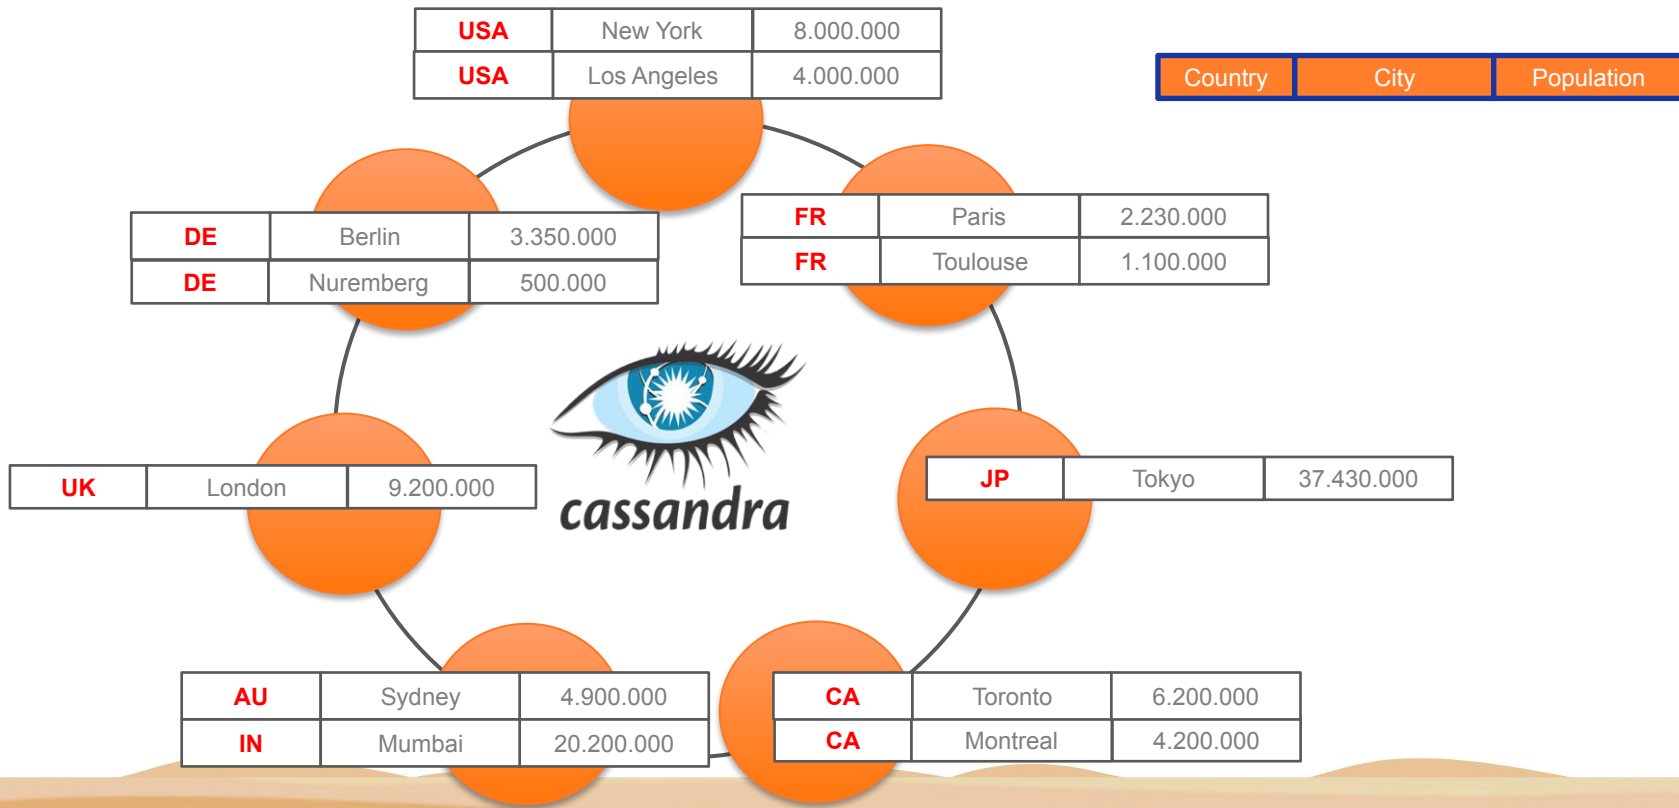
Data is Distributed



Data is Distributed



Data is Distributed






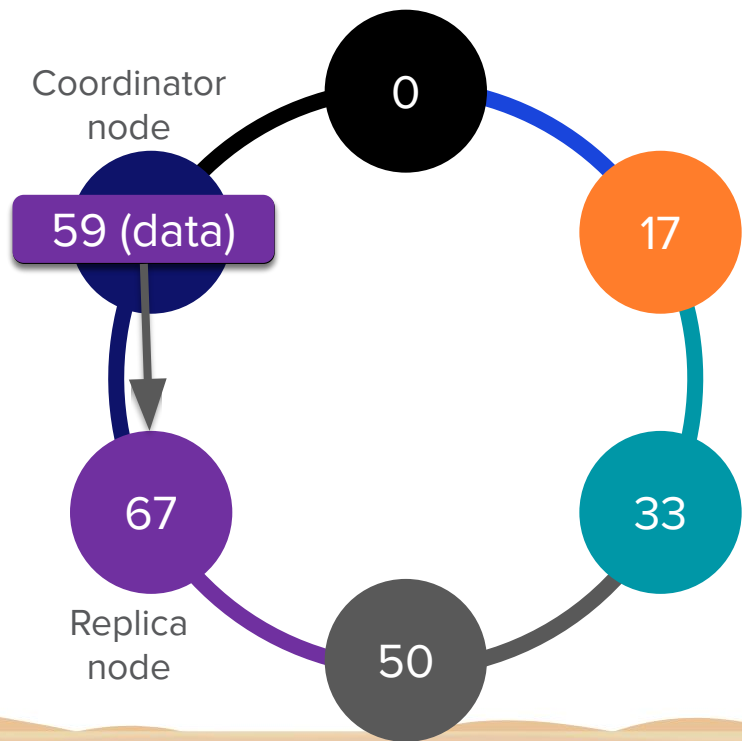
Replication



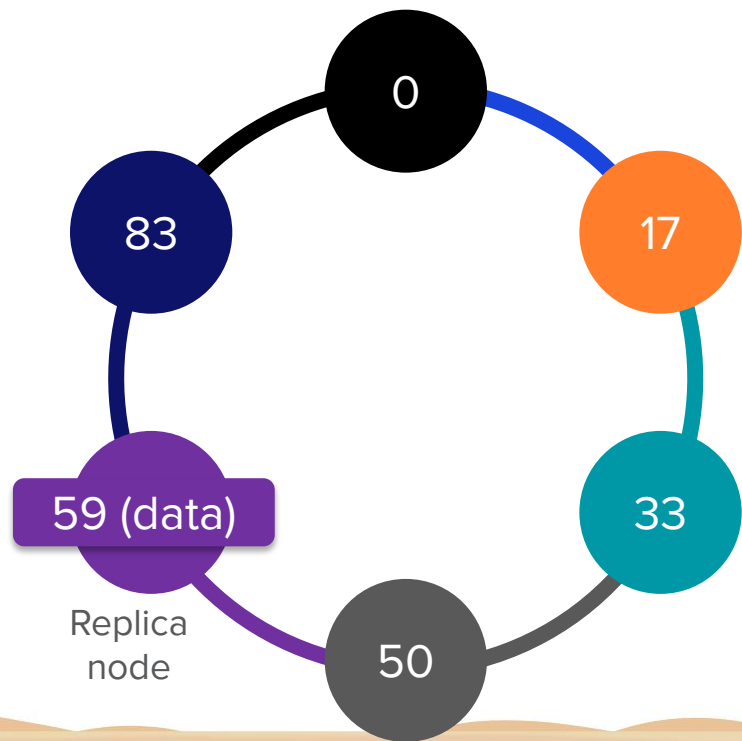
How the Ring Works



How the Ring Works




How the Ring Works




Replication within the Ring

RF = 2




Replication within the Ring

RF = 2




Replication within the Ring

RF = 2




Replication within the Ring

RF = 3




Replication within the Ring

RF = 3




Replication within the Ring

RF = 3



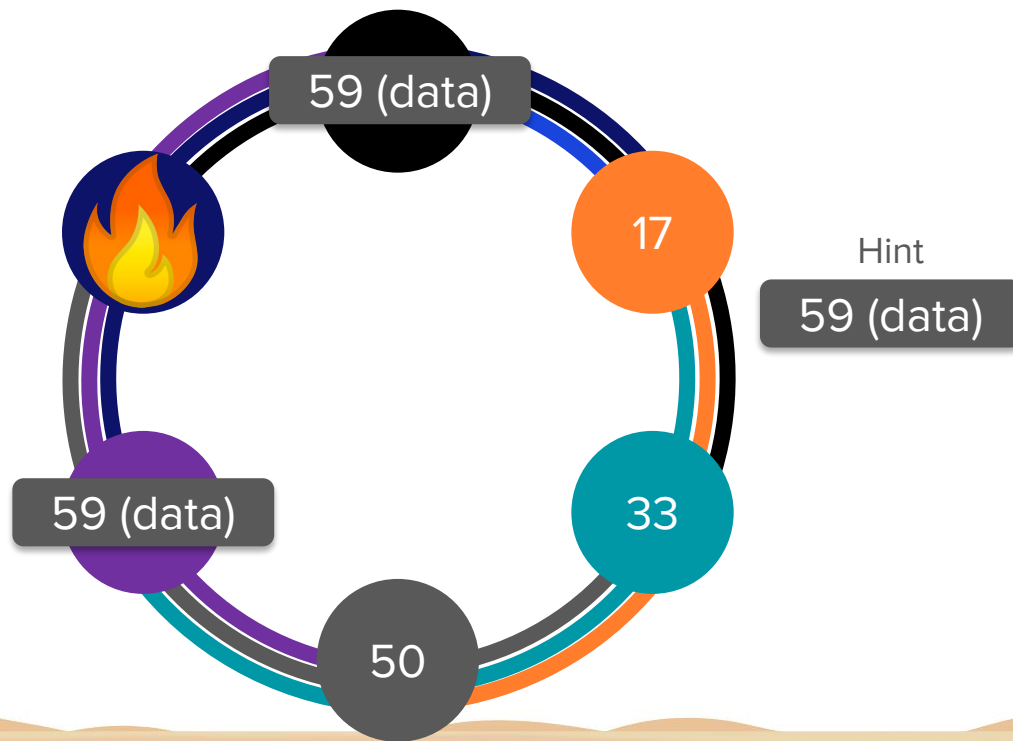
Node Failure

RF = 3

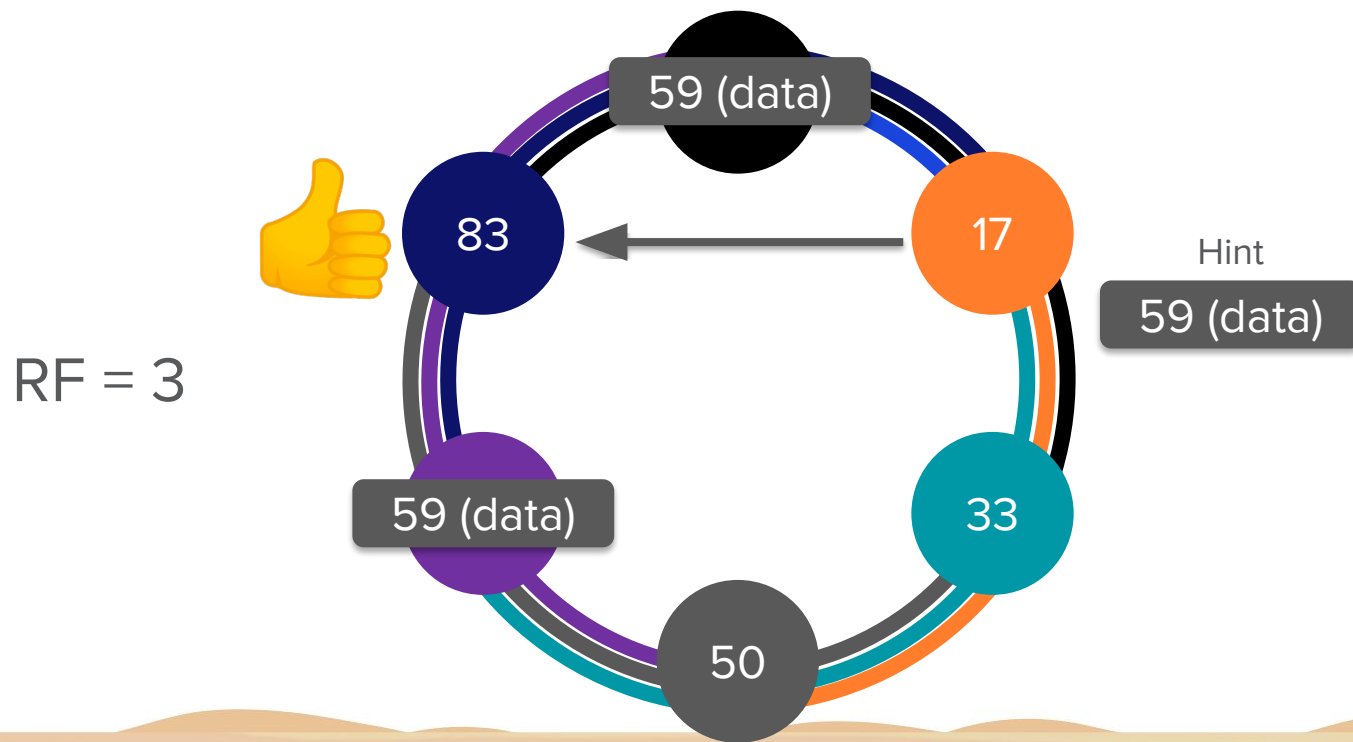


Node Failure

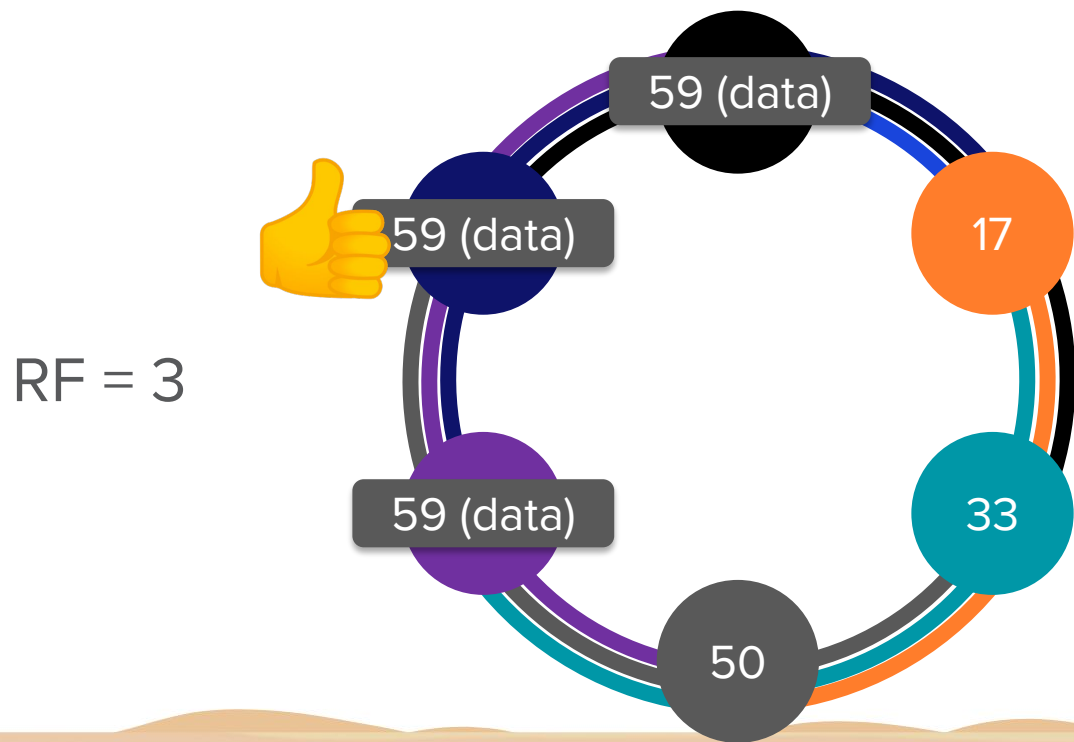
RF = 3



Node Failure




Node Failure – Recovered!




Data Distributed Everywhere

- Geographic Distribution


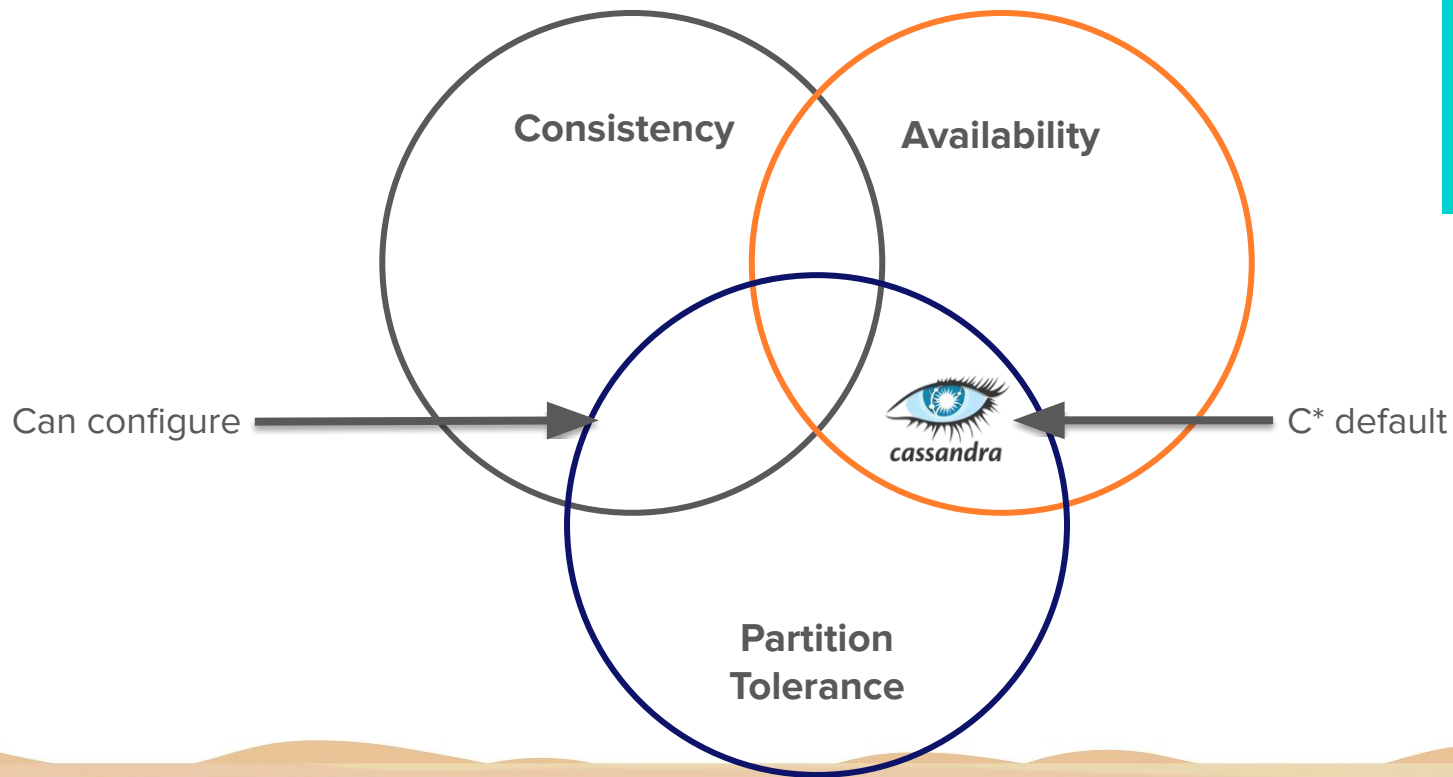


- Hybrid-Cloud and Multi-Cloud




Consistency


CAP Theorem




Consistency Levels



Consistency Levels




Consistency Levels




Immediate Consistency

$$CL_{\text{read}} + CL_{\text{write}} > RF$$

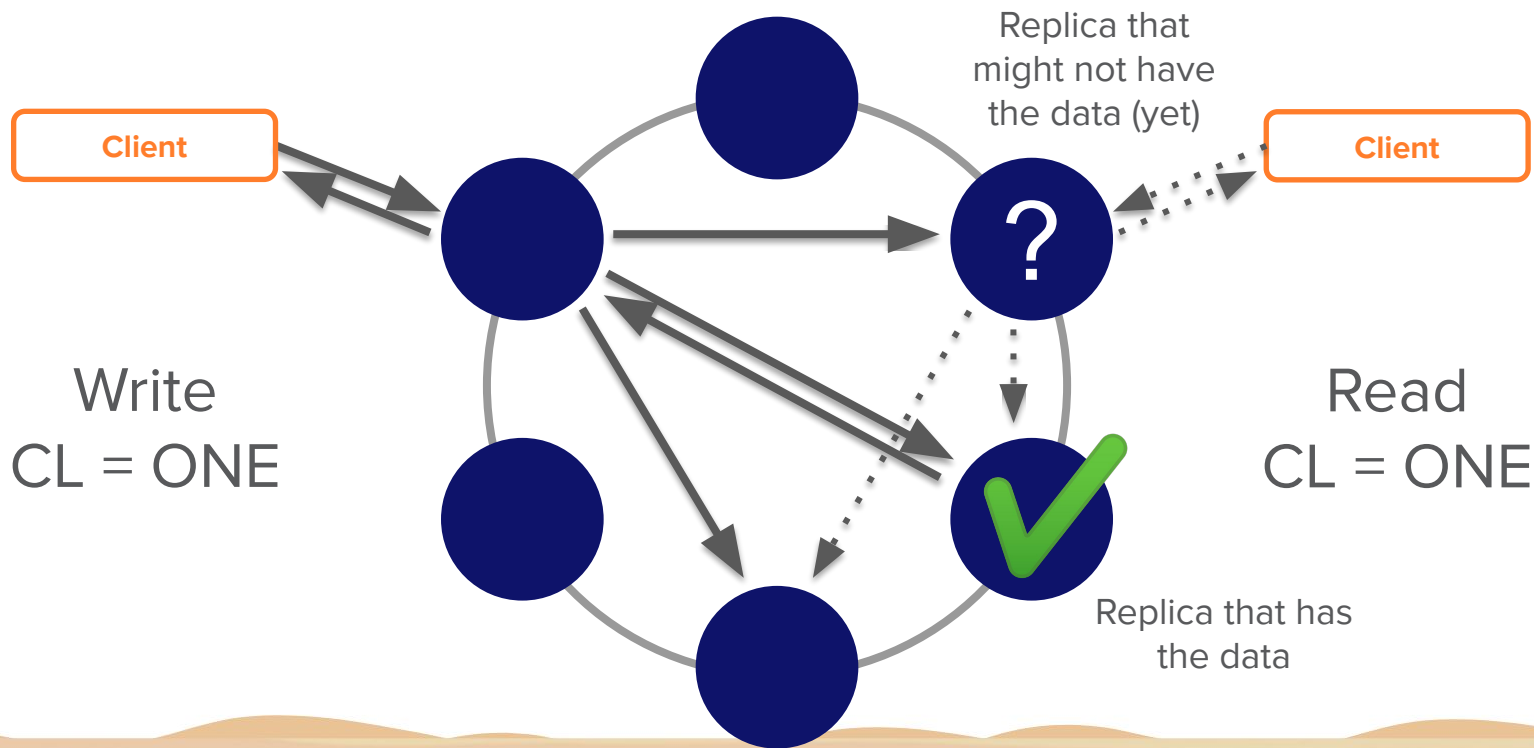
Immediate Consistency – One Way



Immediate Consistency – A Better Way



Weak(er) Consistency



Understanding Use Cases

Scalability

High Throughput
High Volume



Heavy Writes
Heavy Reads



Event Streaming	Log Analytics
Internet of Things	Other Time Series

Availability

Mission-Critical



No Data Loss
Always-on



Caching	Pricing
Market Data	Inventory

Distributed

Global Presence
Workload Mobility



Compliance / GDPR



Banking	Retail
Tracking / Logistics	Customer Experience

Cloud-native

Modern Cloud Applications



API Layer	Hybrid-cloud
Enterprise Data Layer	Multi-cloud

Developer Workshop Series Week 1



**What we will
cover:**

- Bootstrapping
- Apache Cassandra™ Why, What & When
- Read and Write path
- Uber High Level Data Modeling
- What's NEXT?

Write path



Write Path

RAM

DISK

Write Path

RAM



1	Dev Awesome	TX	Houston
---	-------------	----	---------

DISK

Write Path

MemTable

RAM

1	Dev Awesome	TX	Houston
---	-------------	----	---------

Commit
Log

DISK

1	Dev Awesome	TX	Houston
---	-------------	----	---------

Write Path

1	Dev Awesome	TX	Houston
---	-------------	----	---------

MemTable

RAM

ACK

Commit
Log

DISK

1	Dev Awesome	TX	Houston
---	-------------	----	---------

Write Path

1	Dev Awesome	TX	Houston
---	-------------	----	---------

MemTable

RAM



2	Come To Cassandra	TX	Dallas
---	-------------------	----	--------

Commit
Log

DISK

1	Dev Awesome	TX	Houston
---	-------------	----	---------

Write Path

MemTable

RAM

2	Come To Cassandra	TX	Dallas
1	Dev Awesome	TX	Houston

ACK

Commit
Log

DISK

1	Dev Awesome	TX	Houston
2	Come To Cassandra	TX	Dallas

Write Path

MemTable

RAM

2	Come To Cassandra	TX	Dallas
1	Dev Awesome	TX	Houston



Commit
Log

DISK

1	Dev Awesome	TX	Houston
2	Come To Cassandra	TX	Dallas

3	Lone Node	TX	Snyder
---	-----------	----	--------

Write Path

MemTable

RAM

2	Come To Cassandra	TX	Dallas
1	Dev Awesome	TX	Houston
3	Lone Node	TX	Snyder

ACK

Commit
Log

DISK

1	Dev Awesome	TX	Houston
2	Come To Cassandra	TX	Dallas
3	Lone Node	TX	Snyder

Write Path

MemTable

RAM

2	Come To Cassandra	TX	Dallas
1	Dev Awesome	TX	Houston
3	Lone Node	TX	Snyder



Commit
Log

DISK

4	IgotUr Data	TX	Austin
---	-------------	----	--------

1	Dev Awesome	TX	Houston
2	Come To Cassandra	TX	Dallas
3	Lone Node	TX	Snyder

Write Path

MemTable

RAM

4	IgotUr Data	TX	Austin
2	Come To Cassandra	TX	Dallas
1	Dev Awesome	TX	Houston
3	Lone Node	TX	Snyder

ACK

Commit
Log

DISK

1	Dev Awesome	TX	Houston
2	Come To Cassandra	TX	Dallas
3	Lone Node	TX	Snyder
4	IgotUr Data	TX	Austin

Write Path

MemTable

RAM

4	IgotUr Data	TX	Austin
2	Come To Cassandra	TX	Dallas
1	Dev Awesome	TX	Houston
3	Lone Node	TX	Snyder



Commit
Log

DISK

5	Lone Star	TX	El Paso
---	-----------	----	---------

1	Dev Awesome	TX	Houston
2	Come To Cassandra	TX	Dallas
3	Lone Node	TX	Snyder
4	IgotUr Data	TX	Austin

Write Path

MemTable

RAM

4	IgotUr Data	TX	Austin
2	Come To Cassandra	TX	Dallas
5	Lone Star	TX	El Paso
1	Dev Awesome	TX	Houston
3	Lone Node	TX	Snyder

ACK

Commit
Log

DISK

1	Dev Awesome	TX	Houston
2	Come To Cassandra	TX	Dallas
3	Lone Node	TX	Snyder
4	IgotUr Data	TX	Austin
5	Lone Star	TX	El Paso

Write Path

MemTable

RAM

4	IgotUr Data	TX	Austin
2	Come To Cassandra	TX	Dallas
5	Lone Star	TX	El Paso
1	Dev Awesome	TX	Houston
3	Lone Node	TX	Snyder

FLUSH

SSTABLE

Commit
Log

DISK

1	Dev Awesome	TX	Houston
2	Come To Cassandra	TX	Dallas
3	Lone Node	TX	Snyder
4	IgotUr Data	TX	Austin
5	Lone Star	TX	El Paso

4	IgotUr Data	TX	Austin
2	Come To Cassandra	TX	Dallas
5	Lone Star	TX	El Paso
1	Dev Awesome	TX	Houston
3	Lone Node	TX	Snyder

Write Path

MemTable

RAM

Commit
Log

DISK

1	Dev Awesome	TX	Houston
2	Come To Cassandra		Dallas
3	Lone Node	TX	Snyder
4	IgotUr Data		Austin
5	Lone Star	TX	El Paso

SSTABLE

4	IgotUr Data	TX	Austin
2	Come To Cassandra	TX	Dallas
5	Lone Star	TX	El Paso
1	Dev Awesome	TX	Houston
3	Lone Node	TX	Snyder

Write Path

MemTable

RAM

DISK

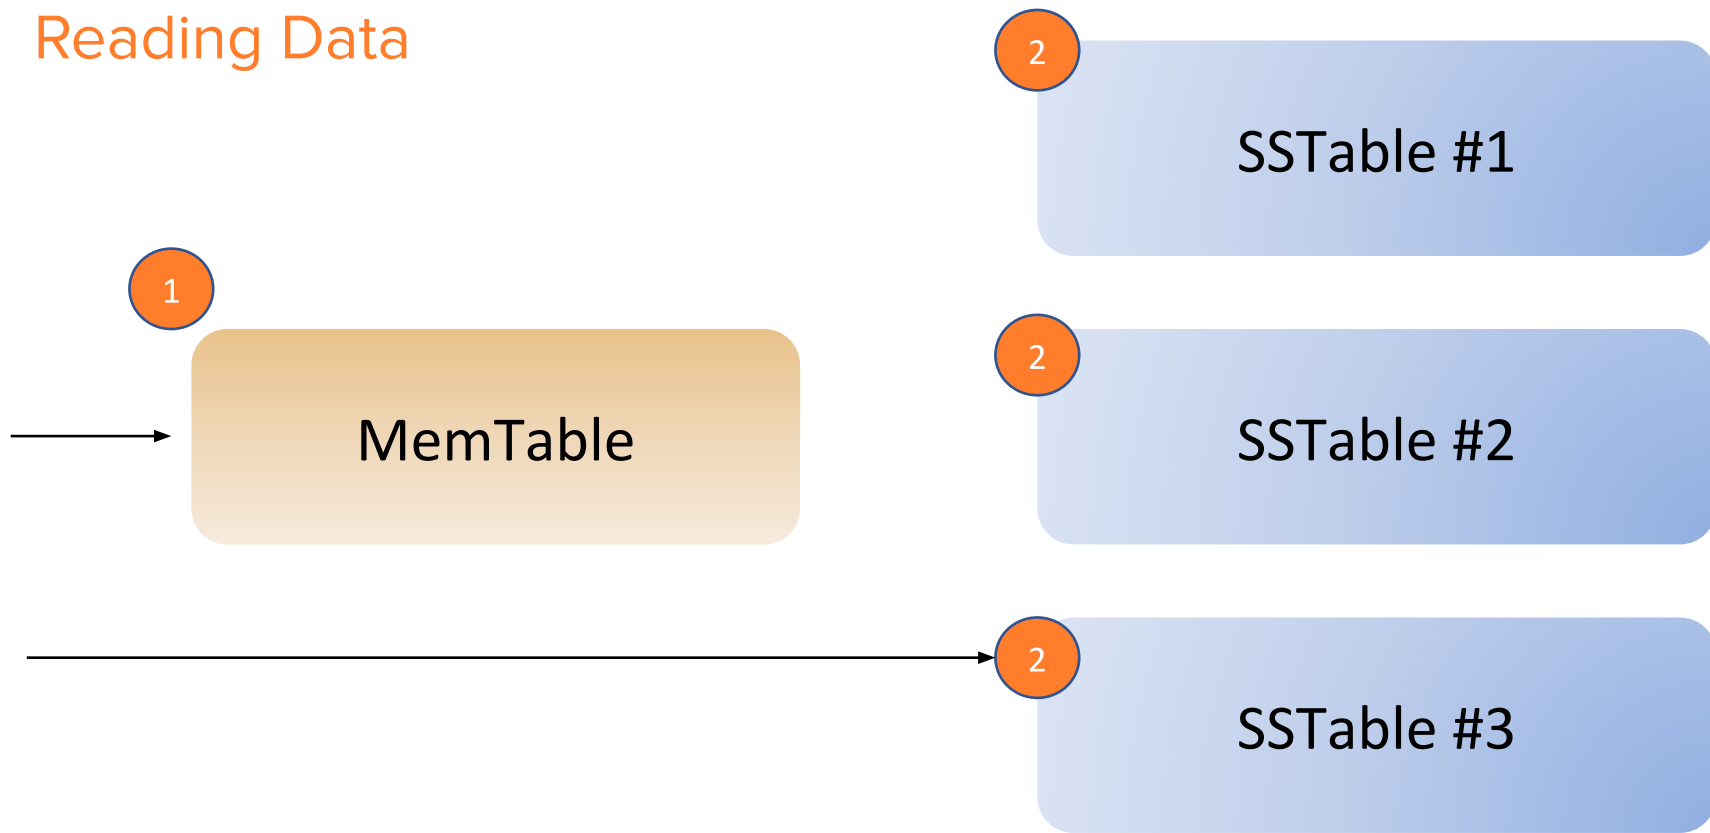
SSTABLE
(IMMUTABLE)

4	IgotUr Data	TX	Austin
2	Come To Cassandra	TX	Dallas
5	Lone Star	TX	El Paso
1	Dev Awesome	TX	Houston
3	Lone Node	TX	Snyder

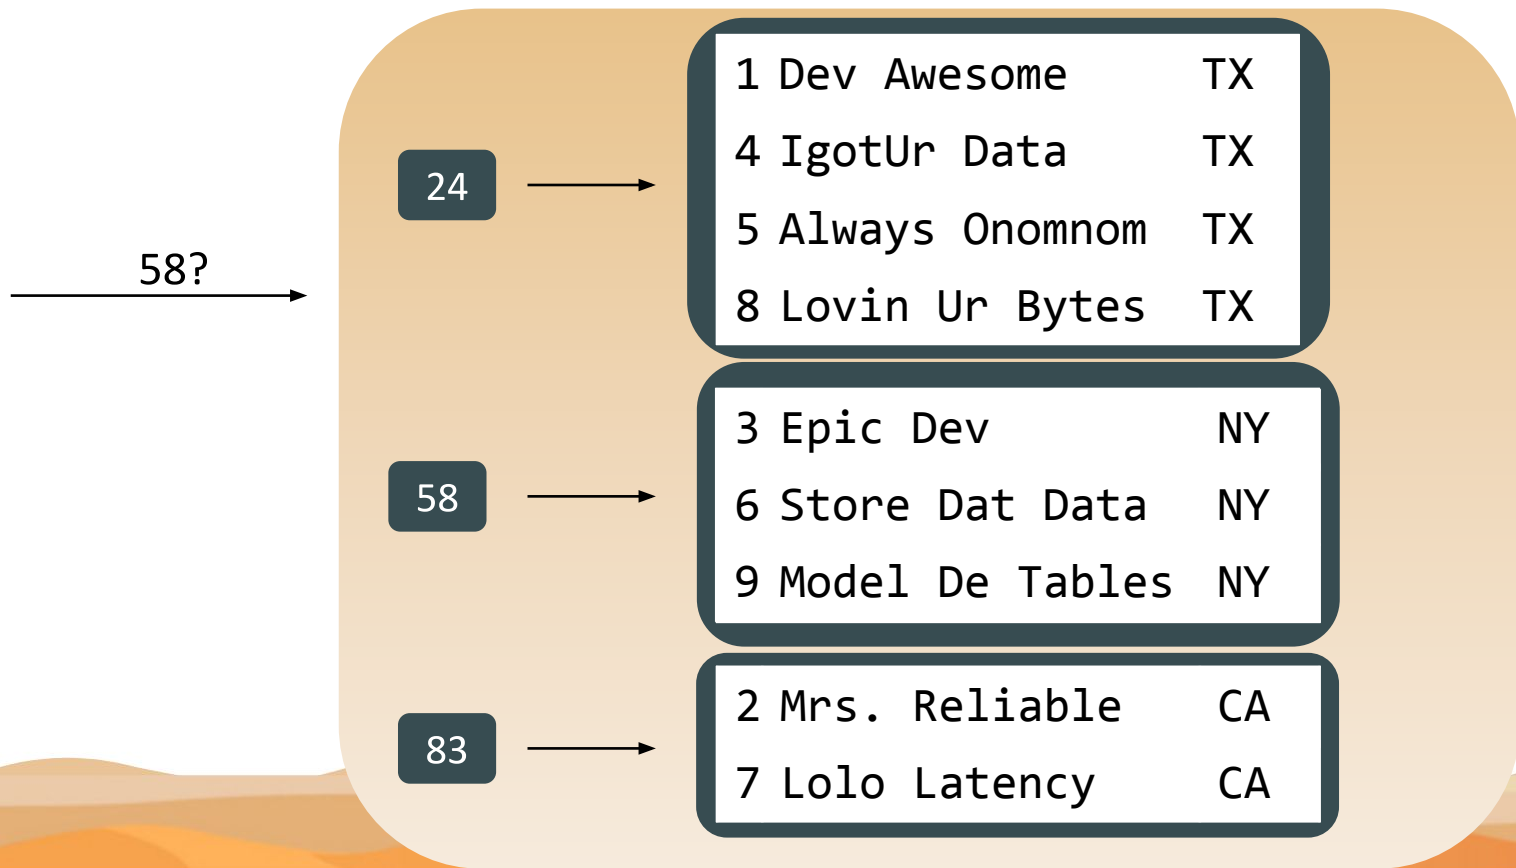
Read path



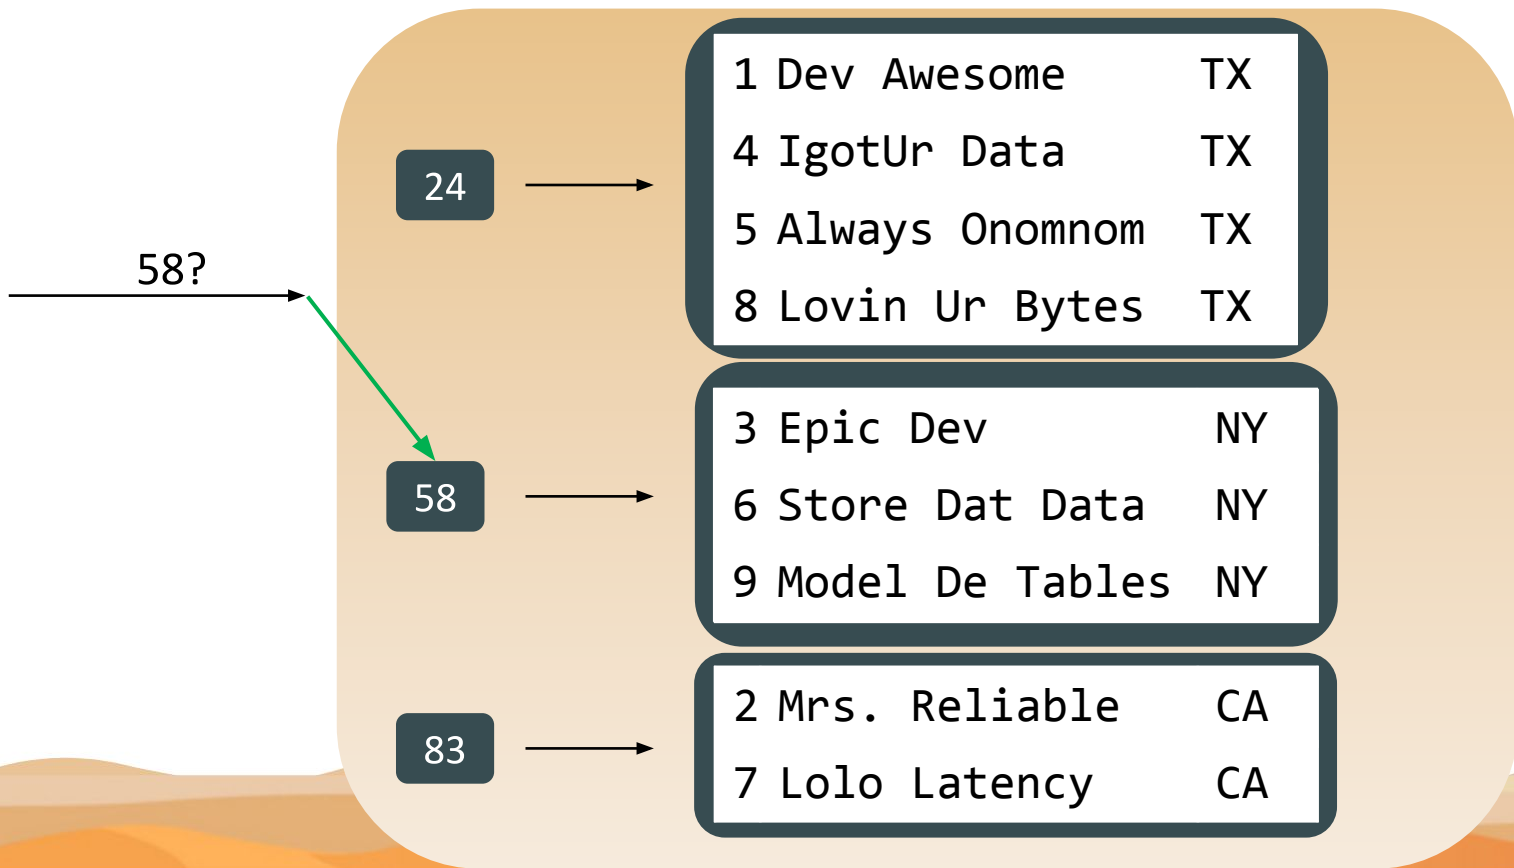
Reading Data




Reading a MemTable



Reading a MemTable




Reading a SSTable



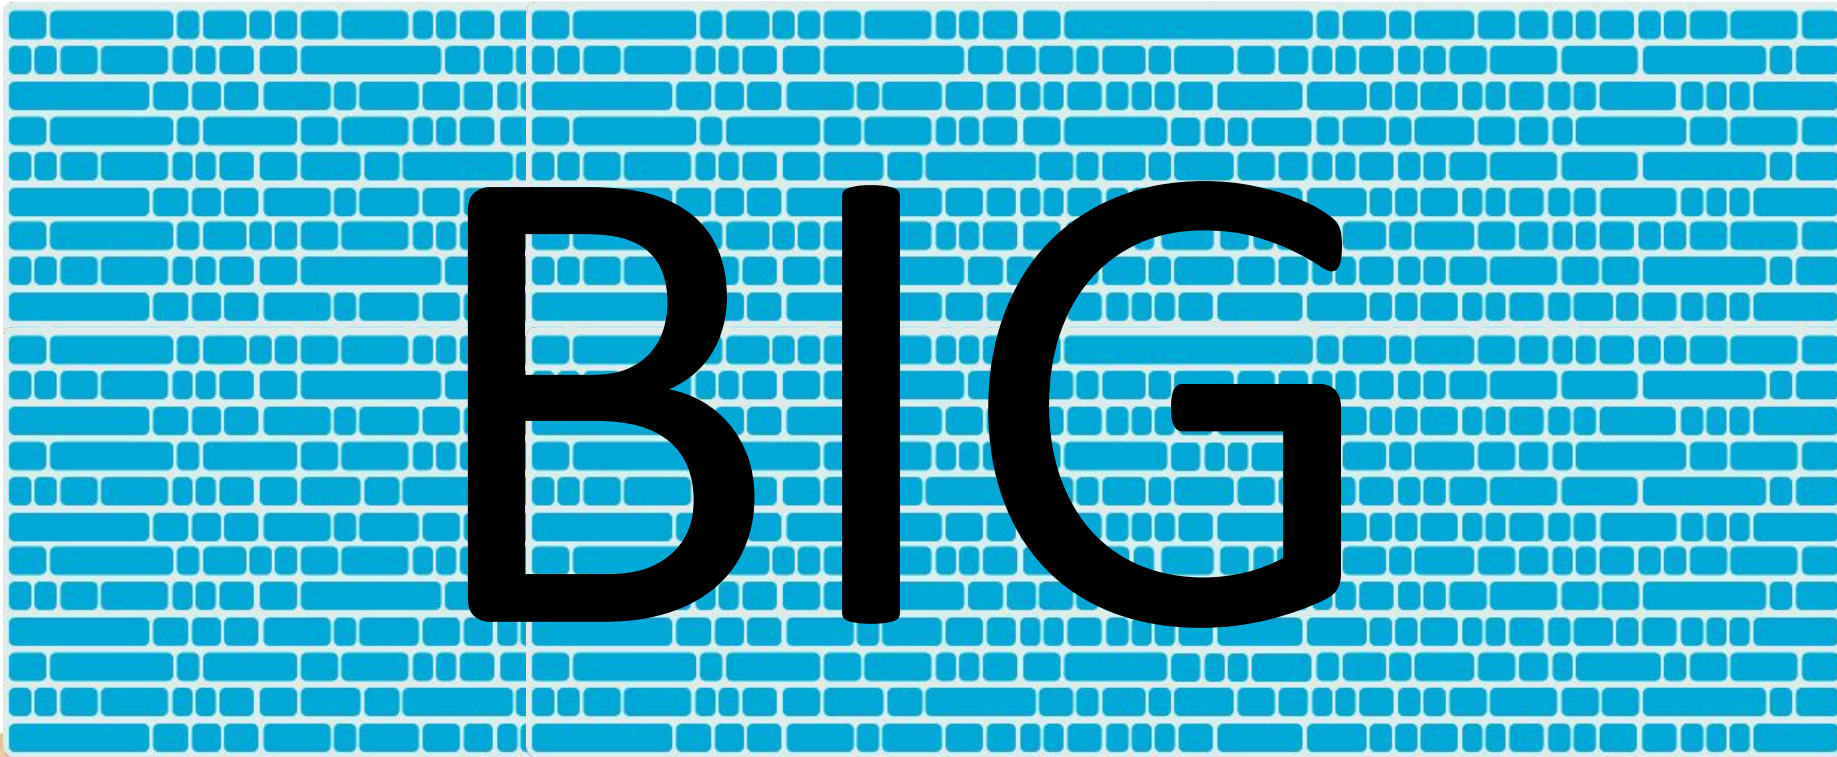
- ❖ A SSTable holds ordered partitions
- ❖ A partition can be split in multiple SSTables
- ❖ We can mark offset of each partition

Partition Index (on Disk)

Token	Byte Offset
7	0
13	1,120
18	3,528
21	4,392
36	6,224
58	7,192
92	9,120



...but a SSTABLE is ...



BIG

Partition Summary (RAM)

Token	Byte Offset
	0
	24
56 - 100	40


Token	Byte Offset
7	0
13	1,120
18	3,528
21	4,392
36	6,224
58	7,192
92	9,120




Key Cache (RAM)

Token	Byte Offset
	0
	24
	40

Token	Byte Offset
7	0
13	1,120
18	3,528
21	4,392
36	6,224
58	7,192
92	9,120



Bloom Filter



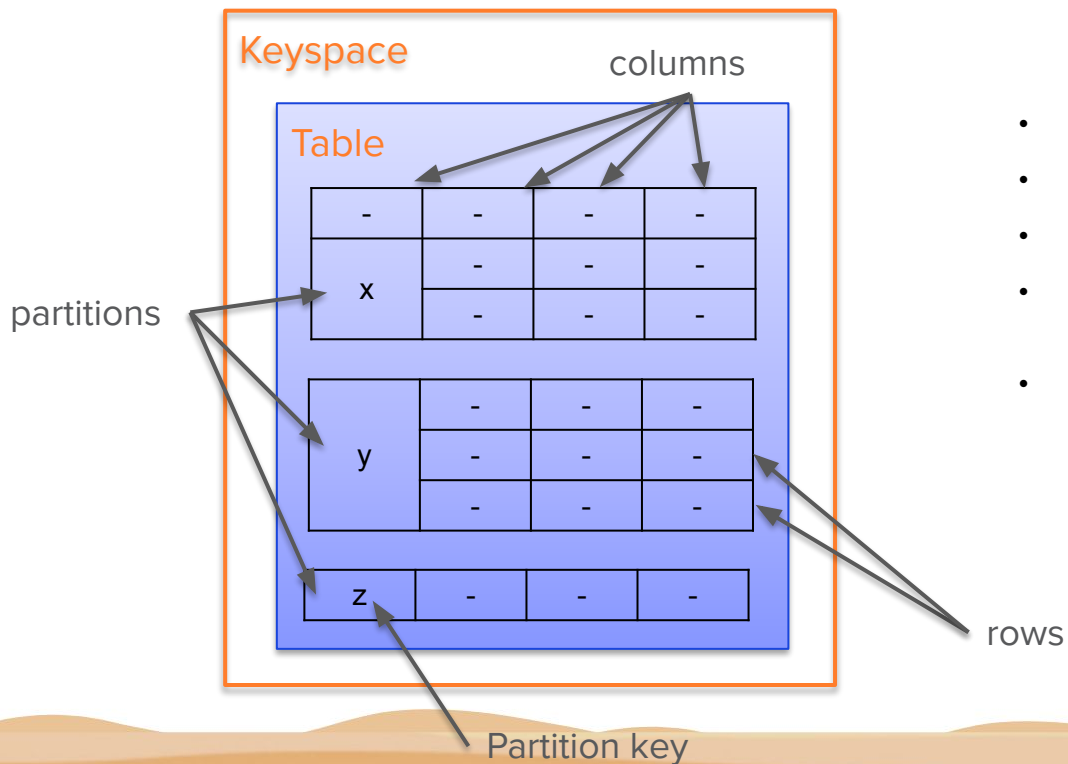
Developer Workshop Series Week 1



**What we will
cover:**

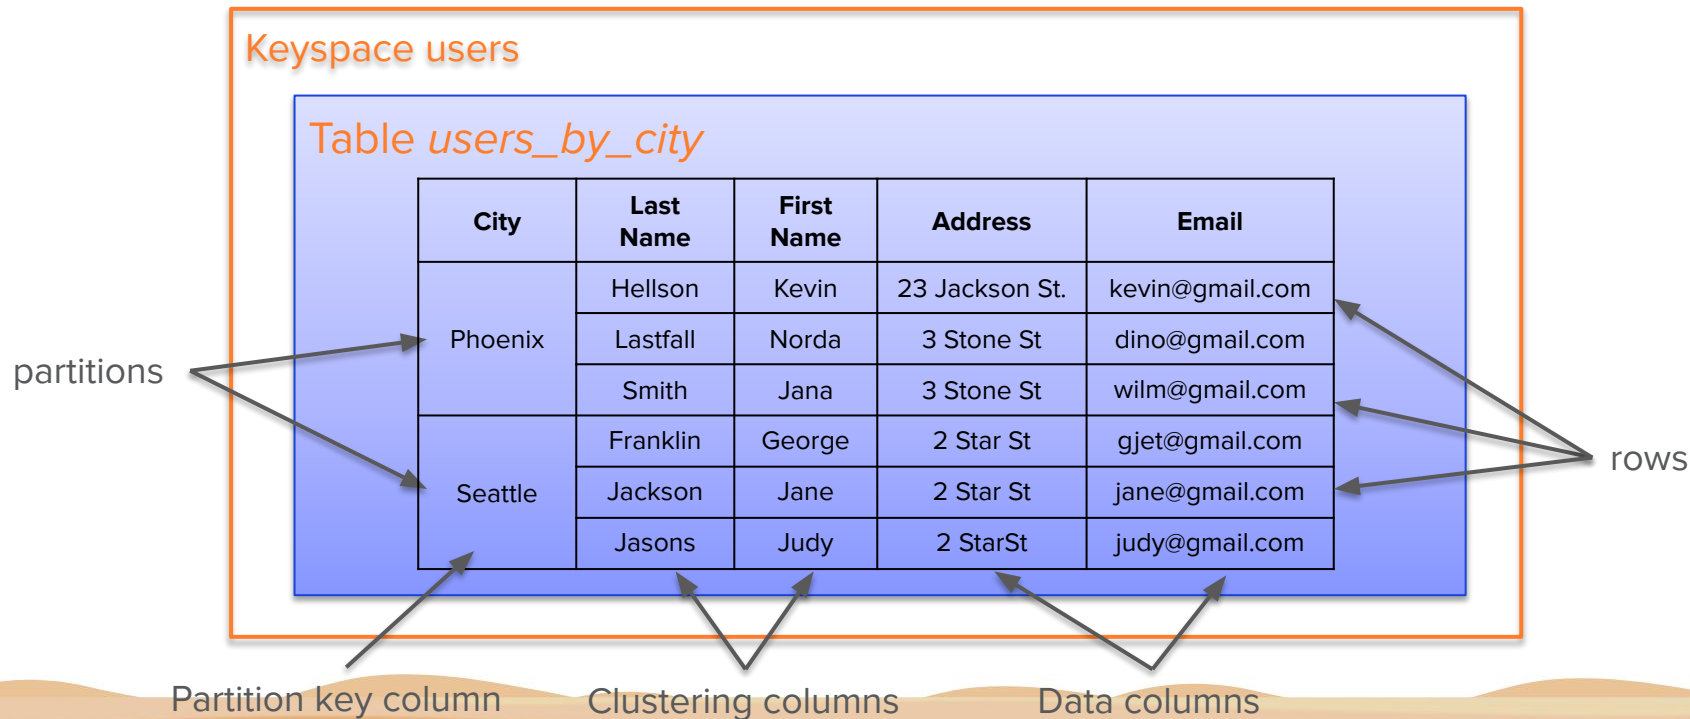
- Bootstrapping
- Apache Cassandra™ Why, What & When
- Read and Write path
- Uber High Level Data Modeling
- What's NEXT?

Cassandra Structure - Partition



- Tabular data model, with one twist
- *Keyspaces* contain *tables*
- *Tables* are organized in *rows* and *columns*
- Groups of related rows called *partitions* are stored together on the same node (or nodes)
- Each row contains a *partition key*
 - One or more columns that are hashed to determine which node(s) store that data

Example Data – Users organized by city



Tables Hold Many Partitions

City	Last Name	First Name	Address	Email
Phoenix	Hellson	Kevin	23 Jackson St.	kevin@gmail.com
	Lastfall	Norda	3 Stone St	dino@gmail.com
	Smith	Jana	3 Stone St	wilm@gmail.com

Table *users_by_city*

Tables Hold Many Partitions

City	Last Name	First Name	Address	Email
Seattle	Franklin	George	2 Star St	gjet@gmail.com
	Jackson	Jane	2 Star St	jane@gmail.com
	Jasons	Judy	2 StarSt	judy@gmail.com

Table *users_by_city*

City	Last Name	First Name	Address	Email
Phoenix	----	----	----	----
	----	----	----	----
	----	----	----	----

Tables Hold Many Partitions

City	Last Name	First Name	Address	Email
Charlotte	Azrael	Chris	5 Blue St	chris@gmail.com
	Stilson	Brainy	7 Azure Ln	brain@gmail.com
	Smith	Cristina	4 Teal Cir	clu@gmail.com
	Sage	Grant	9 Royal St	grant@gmail.com
	Seterson	Peter	2 Navy Ct	peter@gmail.com

Table *users_by_city*

City	Last Name	First Name	Address	Email
Phoenix	----	----	----	----
	----	----	----	----
	----	----	----	----
Seattle	----	----	----	----
	----	----	----	----
	----	----	----	----

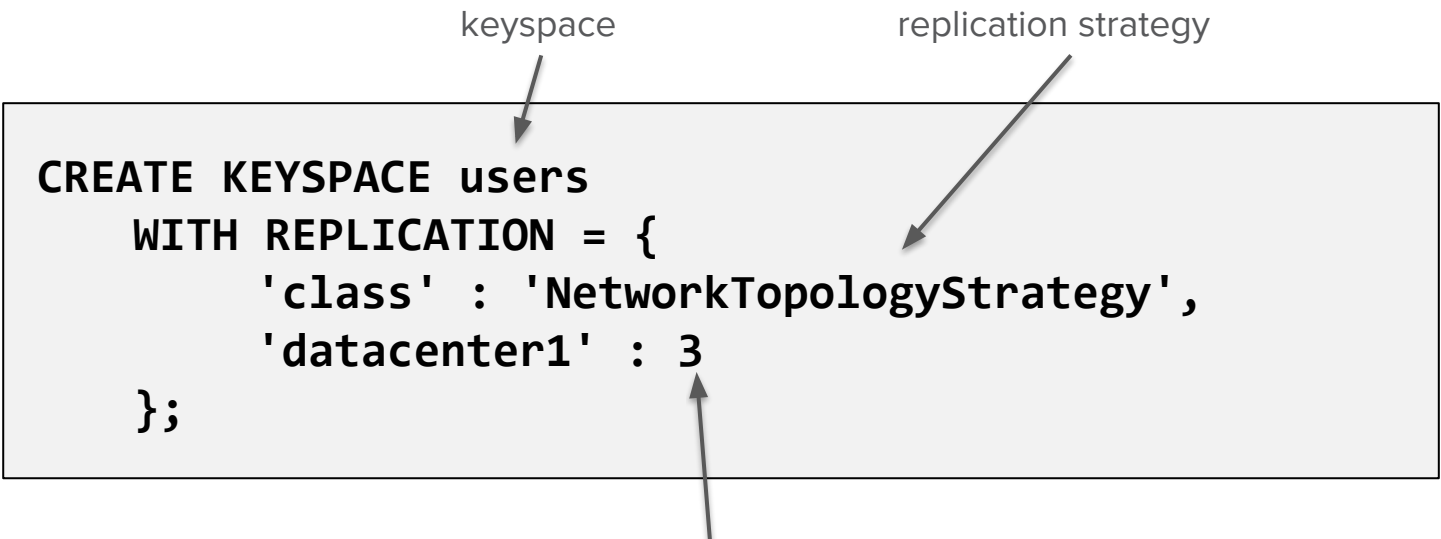
Tables Hold Many Partitions

Table *users_by_city*

City	Last Name	First Name	Address	Email
Phoenix	----	----	----	----
	----	----	----	----
	----	----	----	----
Seattle	----	----	----	----
	----	----	----	----
	----	----	----	----
Charlotte	----	----	----	----
	----	----	----	----
	----	----	----	----
	----	----	----	----
	----	----	----	----

Creating a Keyspace in CQL

```
CREATE KEYSPACE users
  WITH REPLICATION = {
    'class' : 'NetworkTopologyStrategy',
    'datacenter1' : 3
  };
```



Replication factor by data center

Creating a Table in CQL

